

ITSC 2214 - Data Structures and Algorithms - Summer 2023

2023-05-22

Table of contents

Preface	9
I On Computers and Computing	10
1 Background on Computer Science and Data Structures	11
1.1 Computing Capabilities	11
1.2 What is Computer Science then?	14
1.3 Data Structures and Algorithms	14
1.4 Review	14
2 Modern General-Purpose, Programmable Computers	15
2.1 The Working of a Processor	15
2.2 Memory	17
2.3 Programs	17
2.4 The Java Compiler	17
2.5 Operating Systems	18
II Algorithmic Analysis	19
3 Algorithmic Analysis	20
3.1 Problems, Algorithms, and Programs	20
3.1.1 Problems	20
3.1.2 Algorithms	20
3.1.3 Programs and Their Building Blocks	21
3.2 Comparing the Performance of Programs	22
3.3 Analyzing Algorithms	23
3.3.1 Predicting Execution Time	23
3.3.2 Impact of Input Size	23
3.3.3 Iterations and Input Size	24
3.3.4 Gauging Relative Execution Time	24
3.4 Algorithm Complexities	24
3.4.1 Time Complexity	24
3.4.2 Space Complexity	25
3.4.3 Conditional Statements and Complexity	25

3.5	Common Big-O Complexities	26
3.5.1	Constant Complexity and Growth	26
3.5.2	Linear Complexity and Growth	26
3.5.3	Quadratic Complexity and Growth	27
3.5.4	Exponential Functions and Growth	27
3.5.5	Logarithmic Functions and Growth	27
3.5.6	Examples	28
3.6	Growth Rate	29
3.7	Summary / Review	29
III	Unit Testing	31
4	Unit Testing and Test-Driven Development	32
4.1	Background	32
4.2	The Power of Unit Testing	32
4.3	A Basic Approach to Unit Testing	33
4.4	Recognizing the Inefficiencies and Redundancies	34
4.4.1	Redundancy	34
4.4.2	Lack of Automation	35
4.5	A slightly more sophisticated approach to Unit Testing	36
4.6	Automating Test Execution	37
4.7	Summary	37
5	Writing JUnit Tests and Test-Driven Development	39
5.1	Introduction to JUnit	39
5.2	Useful JUnit Assertions	40
5.2.1	<code>assertEquals</code>	40
5.2.2	<code>assertTrue</code> and <code>assertFalse</code>	40
5.2.3	<code>assertNotNull</code> and <code>assertNull</code>	41
5.3	Setting Up JUnit	41
5.4	Utilizing <code>@Test</code> and <code>@BeforeEach</code> Annotations	42
5.5	Interpreting JUnit Test Runner Output	42
5.6	Conclusion	44
5.7	Test-Driven Development	44
5.7.1	The Motivation Behind Test-Driven Development	45
6	Introduction to Mutation Testing	46
6.1	Unit Testing, a Recap	46
6.2	Writing <i>Correct</i> Unit Tests	47
6.3	Mutation Testing: A Solution to Test the Correctness of Your Unit Tests	47
6.4	Mutation Testing in Practice	47
6.4.1	What are these mutators?	48

6.5	Understanding CodingRooms Feedback	49
6.5.1	Understanding Mutation Test Feedback	49
6.5.2	Understanding the Grading Overview	50
7	List of Mutators	52
7.1	Conditionals Boundary Mutator (CONDITIONALS_BOUNDARY)	52
7.2	Increments Mutator (INCREMENTS)	52
7.3	Invert Negatives Mutator (INVERT_NEGS)	53
7.4	Math Mutator (MATH)	53
7.5	Negate Conditionals Mutator (NEGATE_CONDITIONALS)	55
7.6	Return Values Mutator (RETURN_VALS)	56
7.7	Void Method Call Mutator (VOID_METHOD_CALLS)	58
7.8	Empty returns Mutator (EMPTY RETURNS)	58
7.9	False returns Mutator (FALSE_RETURNS)	59
7.10	True returns Mutator (TRUE_RETURNS)	59
7.11	Null returns Mutator (NULL_RETURNS)	59
7.12	Primitive returns Mutator (PRIMITIVE_RETURNS)	60
7.13	Constructor Call Mutator (CONSTRUCTOR_CALLS)	60
IV	Iterative Sorting and Searching	61
8	Introduction to Sorting	62
8.1	Understanding the Sorting Problem	62
8.2	Developing a Sorting Algorithm from Scratch	62
8.3	Starting Simple: Sorting an Array of Size 1	62
8.4	Sorting an Array of Size 2	63
8.5	Expanding the Problem: Sorting an Array of Size 3	63
8.6	Scaling Up: Sorting an Array of Size 4	64
8.7	Recognizing the Pattern	64
8.8	Generalizing the Algorithm	65
9	Comparable and Comparator	66
9.1	Sorting Complex Objects	66
9.2	Defining a Comparison Method	66
9.3	Implementing Comparable	66
9.4	Comparable in the Standard Library	67
9.5	Sorting Flexibility with Comparator	68
9.6	Understanding Natural and Total Orderings	68
10	Bubble, Selection, and Insertion Sort	70
10.1	Bubble Sort	70
10.1.1	Pseudocode	70

10.1.2 An Example	71
10.1.3 Lab Instructions	72
10.2 Selection Sort	73
10.2.1 Introduction	73
10.2.2 Pseudocode	73
10.2.3 An Example	74
10.2.4 Instructions	75
10.3 Insertion Sort	75
10.3.1 Introduction	75
10.3.2 Pseudocode	75
10.3.3 An Example	76
10.3.4 Instructions	78
11 Linear Search	79
11.1 Pseudocode	79
11.2 An Example	80
11.3 Instructions	80
V Sequential Collections of Data	82
12 Collections of Data	83
12.1 (Active) Memory in Computers	83
12.2 The Concept of Data Collections	83
12.3 Types of Collections: Fixed and Dynamic	84
12.4 Memory Layout of Collections	85
12.5 Tackling Dynamic Collections	85
12.6 Memory Layout and Program Performance	85
12.7 Summary and Conclusion	87
13 Operations on Lists	88
13.1 Adding Data to a List	88
13.1.1 Appending an Element	89
13.1.2 Prepending an Element	89
13.1.3 Inserting an Element at a Specific Index	89
13.1.4 Adding All Elements from Another List	89
13.2 The Role of Abstract Data Types (ADTs)	90
13.2.1 Summary of Addition Operations	90
13.3 Removing data from our List	91
13.4 Removing Data from a List	91
13.4.1 Removing an Element at a Specific Index	91
13.4.2 Removing the First Occurrence of an Element	91
13.4.3 Removing All Elements from Another List	92

13.4.4	Removing All Elements from the List	92
13.4.5	Summary of Removal Operations	92
13.5	Searching for data in our List	93
13.6	Searching in a List	93
13.6.1	Checking If an Element Exists in the List	93
13.6.2	Checking If All Elements of Another Collection Exist in the List	93
13.6.3	Finding the Index of an Element	93
13.6.4	Finding the Last Index of an Element	94
13.6.5	Summary of Search Operations	94
13.7	Miscellaneous Operations on a List	94
13.7.1	Accessing an Element	95
13.7.2	Modifying an Element	95
13.7.3	Determining the Size of the List	95
13.7.4	Converting the List to an Array	95
13.7.5	Summary of Miscellaneous Operations	95
13.8	Summary and Conclusion	96
14	Implementing Array-backed Lists	98
15	Implementing Linked Lists	99
15.1	Understanding References in Java	99
15.2	Memory Allocation for Java Objects	100
15.3	6.3 Objects Interacting Through References	100
15.4	Implementing Linked Lists in Java	103
15.5	6.5 Adding to a Singly Linked List	106
15.6	Searching a Singly Linked List	109
15.7	Removing from a Singly Linked List	111
15.7.1	Case 1: Empty Linked List	111
15.7.2	Case 2: Linked List with One Element	111
15.7.3	Case 3: Linked List with More Than One Element	112
15.7.4	Moving to <code>remove(int i)</code>	112
15.7.5	Summary and Conclusion	114
VI	Graphs	115
16	The Graph Data Structure	116
16.1	Background and Motivation	116
16.2	Introduction	119
16.3	Graph Terminology	119
16.3.1	Basic Terms and Properties	119
16.3.2	Graph Notation	129
16.3.3	Special Types of Graphs	129

16.4	Graph Representation	132
16.4.1	Adjacency List	132
16.4.2	Adjacency Matrix	133
16.4.3	Converting Between Representations	135
16.5	Graph Traversal	135
16.5.1	Breadth-First Search (BFS)	136
16.5.2	Depth-First Search (DFS)	138
	16.5.3 Applications and Variations of BFS and DFS	140
VII	Hashing, HashMaps and HashTables	142
17	Hashing, Hash Tables, and Hash Maps	143
17.1	Background and Motivation	143
17.1.1	Indexing	143
17.1.2	Limitations of Indexing	143
17.1.3	Mapping Strings to Numbers	143
17.1.4	Hashing: A Better Solution	144
17.2	Hash Functions	145
17.2.1	Introduction	145
17.2.2	Hashing	145
17.2.3	Simple Hash Functions	147
17.2.4	Other Types of Hash Functions	147
	17.2.5 Trade-offs Between Different Hash Functions	149
17.3	Hash Collisions	149
17.4	Chaining	149
17.4.1	Insertion, Search, and Deletion	150
17.4.2	Advantages and Disadvantages of Chaining	150
17.5	Open Addressing	150
17.5.1	Probing Techniques	151
17.5.2	Insertion, Search, and Deletion	151
	17.5.3 Advantages and Disadvantages of Open Addressing	151
17.6	Complexity and Load Factor	152
17.6.1	Time Complexity of Hash Computation	152
17.6.2	Time Complexity of List Traversal	152
17.6.3	Load Factor	152
	17.6.4 Balancing Load Factor and Complexity	153
17.7	Rehashing	153
17.7.1	Why?	153
17.7.2	How?	153
17.8	Hash Tables vs Hash Maps	154
17.9	HashMaps in Java	155
17.10	HashTable in Java	157

17.11HashSets in Java	159
17.12hashCode and equals in Java	161
17.12.1 The hashCode Method	161
17.12.2 Overriding the hashCode Method	161
17.12.3 Using hashCode with Java Collections	162
References	163

Preface

This website contains a set of readings for ITSC 2214 - Data Structures and Algorithms.

Part I

On Computers and Computing

1 Background on Computer Science and Data Structures

Let's break down the term "Computer Science" into its two parts: **computer** and **science**. What do these words mean?

The word **science** means "knowledge" in Latin. Science is a way of finding out things about the world by asking questions, doing experiments, and looking for evidence.

The word **computer** means "to calculate" in Latin. A computer is a machine that can do math problems and store information very fast. The word **computing** means using computers or other machines to solve problems or do tasks.

But did you know that computers were not always machines? Before electronic computers were invented, there were people who did math problems by hand for scientists and engineers. They were called **human computers**. They followed fixed rules and had no authority to deviate from them in any detail. They worked in teams and checked each other's results for accuracy.

1.1 Computing Capabilities

Computers, in the broadest sense, are devices that can perform calculations or manipulate information. Throughout history, humans have invented and used various types of computers, each with increasing capabilities and complexity. Here are some examples of how computing capabilities have evolved over time:

- **Tally stick** - One of the earliest forms of computers, dating back to prehistoric times. A tally stick is a piece of wood or bone with notches carved into it to record numbers or events. The computations it could do were incrementing and retrieving one piece of information. For example, a shepherd could use a tally stick to keep track of his sheep by making a notch for each one.
- **Abacus** - A manual device used for calculations by sliding counters along rods or in grooves. The abacus was invented in ancient times and is still used today in some parts of the world. It could store one set of numbers, add, subtract, multiply, divide, and perform other arithmetic operations to the stored information which can later be retrieved. For example, a merchant could use an abacus to keep track of his transactions and profits.



NASA Dryden Flight Research Center Photo Collection
<http://www.dfrc.nasa.gov/gallery/photo/index.html>
NASA Photo: E49-54 Date: 1949

NACA High Speed Flight Station "Computer Room"

Figure 1.1: For example, in 1959, NASA had a team of women who worked as human computers to help launch rockets into space. They used pencils, paper, and calculators to do complex calculations that machines could not do at that time.

- **Astrolabe** - A sophisticated instrument used for astronomy and navigation by measuring the positions and movements of celestial bodies. The astrolabe was developed in ancient Greece and reached its peak in the Islamic Golden Age. It could perform complex calculations such as determining the time, latitude, longitude, and direction based on the observation of stars and planets. For example, a sailor could use an astrolabe to find his way across the sea by aligning it with the sun or the pole star.
- **Antikythera mechanism** - A mechanical device that simulated the motions of the sun, moon, and planets according to a geocentric model. The Antikythera mechanism was discovered in a shipwreck near the Greek island of Antikythera in 1901. It is estimated to date back to the 2nd century BC and is considered one of the first analog computers. It could predict astronomical phenomena such as eclipses, phases of the moon, and positions of the zodiac signs. For example, a priest could use the Antikythera mechanism to plan religious ceremonies and festivals based on the celestial calendar.
- **Difference engine** - A mechanical calculator that could compute polynomial functions using the method of finite differences. The difference engine was designed by Charles Babbage in the early 19th century but was never fully completed due to technical and financial difficulties. It could generate accurate tables of values for various mathematical functions such as logarithms, trigonometry, and navigation. For example, a mathematician could use the difference engine to check his calculations and avoid errors.
- **Analytical engine** - A proposed mechanical computer that could perform any calculation given a set of instructions or a program. The analytical engine was also designed by Charles Babbage in the mid-19th century but was never built due to his death and lack of funding. It is considered the first general-purpose computer and the precursor of modern computers. It could store data in memory, process data using arithmetic and logical operations, control the flow of execution using conditional branching and looping, and output data using a printer or a punch card. For example, Ada Lovelace, who wrote the first algorithm for the analytical engine, envisioned that it could compose music based on mathematical rules.

Here is a summary of the computing capabilities of some of these devices -

- **Tally stick:** Store and retrieve one piece of data
- **Abacus:** Store and retrieve one piece of data, and perform basic arithmetical operations on them with another operand.
- **System of Gears:** Store and retrieve one piece of data. Each gear can use stored data and scale it up or down (multiply or divide) by a fixed constant determined by the gear ratio.
- **Difference engine:** Can perform complex, but non-programmable computations. Produced tables of input-output pairs.
- **Analytical engine:** Can perform complex, programmable computations. Can store data in memory, process data using arithmetic and logical operations, control the flow of execution using conditional branching and looping, and output data using a printer or a punch card.

1.2 What is Computer Science then?

The science of computers, or Computer Science, seeks to answer fundamental questions like: What are the essential parts of a computer? What can be computed, and what cannot? What determines the ease and speed of computation? This field provides the foundation for understanding the principles that drive computational systems.

1.3 Data Structures and Algorithms

The field of Data Structures and Algorithms expands upon the principles of Computer Science. It determines the efficiency of computation by answering the question: How easily or quickly can something be computed, and what factors influence these metrics? It outlines the necessary building blocks or tools required for programming, and explores the common patterns and problems in programs, offering known ways to enhance their speed. It is this branch of Computer Science that gives us the skills to design, write and analyze the efficiency of our programs. This understanding is crucial to becoming proficient in the broader field of Computer Science.

In essence, the intertwined journey of computers and computing science is a testament to human ingenuity and the relentless pursuit of understanding and harnessing the principles that underlie our world. As we delve deeper into the concepts of data structures and algorithms, we continue to contribute to this exciting journey.

1.4 Review

We talked a little about what computers are, and what capabilities a device needs to have to be able to compute certain types of problems. The next chapter will be about “general purpose, programmable computers”, and how they work.

Here's a fun exercise for you -

- Are analog wristwatches computers? What do they compute?
- Can you design a system of gears that can convert Fahrenheit to Celcius?

2 Modern General-Purpose, Programmable Computers

In this chapter, we will explore the functioning of modern, general-purpose computers. We will approach this topic by drawing parallels with a relatable example from an alternate universe, where two boys, Cory and Jamal, develop a unique system of communication.

In this parallel universe, Cory and Jamal are two teenage neighbors that enjoy playing Scrabble, but there's a hitch. Cory's parents are quite religious and in this universe, there exists a prophecy from their scriptures that a game played by youngsters will one day end the world; and so they disapprove of Cory playing Scrabble.

To get around this, Cory and Jamal come up with a clever plan. Cory happens to have two desk lamps that are visible from Jamal's window. They decide that if Cory's mother isn't home, Cory will switch on the right lamp, and if his father isn't home, he will turn on the left one. If both lamps are on, it essentially signals to Jamal that he can come over to play Scrabble.

Cory also has a sister who is usually at the University, but also at home on some weekends. When she is home, she tends to wake up for late-night snacks around 1 am. And so when Cory's sister is home, they can only play until 1 am. To work around this, they add another signal. If Cory switches on his ceiling fan, it will signal that he can come to play, but only until 1 am.

As we see, each additional piece of information that we need to share requires an additional indicator.

2.1 The Working of a Processor

The way Cory and Jamal communicate parallels how a computer's processor operates. A processor has a set of operations it can perform - like copying, moving, adding, subtracting data, jumping to another instruction, conditional branching, and more.

To tell a computer to perform a certain operation, we need several indicators, just like Cory and Jamal's system. Computers use tiny wires that may or may not have electricity running through them. The presence or absence of electricity indicates what the processor should do.

We refer to the absence as a 0 and the presence of electricity as a 1. Each wire conveying either a 0 or a 1 is said to convey one **bit** of information.

This is a fixed group of wires going into the processor to convey such information on what needs to be done. This grouping is known as "Instruction." The number of bits in an instruction tells us the 'size' of the instruction.

A processor usually has a fixed instruction size (64 for a 64-bit processor, 32 for a 32-bit processor, etc.).

Sometimes even more information is needed. For example, for operations like jumping to a different instruction, the processor needs to know where to jump to. For operations like addition, two numbers are needed. Sometimes, along with the operation, the processor also needs to be instructed about where to store the result. We call these pieces of data on which we perform instructions as **operands**. We often refer to instructions as operations.

We cannot have an extra set of wires for each time we need to share another operand with the processor, so we share this data in sequence.

Say 0010 1010 is the code for the add operation, 0000 0000 is the number 0, and 0001 is the number 1. If we want an 8-bit processor to add these together, we might share the following data through the wires -

```
...
0000 0001 // third cycle - second operand
0000 0000 // second cycle - first operand
0010 1010 // first cycle - add operation
```

There is a "clock" mechanism within the processor to signal the start and end of "cycles". Generally the processor "pipelines" its operation such that it can complete one instruction in each cycle. However, many complex instructions like division or square roots take >20 cycles to complete.

A computer processor can be compared to a complex Rube Goldberg machine, with wires that take in the presence or absence of electricity similar to how many Rube Goldberg machines' mechanism is kicked off by a rolling ball or marbles. The "marbles" of electricity roll through the processor, along each wire, triggering complex chain reactions or side effects that are designed to elicit the intended operation. Consider that a processor, as small as it is, frequently contains billions of transistors - use this information to imagine how complex these "Rube Goldberg machines" are.

2.2 Memory

Let's compare the computer's memory to a vast grid filled with tiny cells. These cells can each store some electricity - and again the presence or absence of electricity conveys a **bit** of information. If we continue the previous analogy, it can be compared to a large grid storing balls (data).

The processor has instructions that allow it to select which cells to "read" and transfer those bits to the processor - to be interpreted either as data or instructions. The processor also has instructions that allow it to write to memory.

It is this interaction of a processor being able to "write" instructions and data to memory, and then conditionally "read" and execute them, and then write new instructions or data that makes a processor "programmable".

What we describe in this section is Active memory or RAM. It is the data in RAM that is readily available for the processor to read and write.

2.3 Programs

Programs are essentially detailed sets of instructions that command a computer to execute specific operations. They can be likened to a recipe that a computer follows to achieve a particular task. Programs dictate what steps the computer must take and in what sequence to reach the desired outcome.

The program is stored on disk in a format that is standard for that operating system, and where the first instruction is stored in this format is always known - say, the first instruction in a "main" section. When a program is executed, it is loaded from the computer's storage into RAM, and then the processor reads and executes the first instruction; and so the execution begins.

The complexity of programs can vary greatly, from a simple one that performs basic arithmetic to an intricate operating system like Windows or Linux that manages every aspect of a computer. However, the core characteristic of all programs is the same: they are sequences of instructions that the computer follows.

2.4 The Java Compiler

As instructions are hard to write directly, we make use of programs called compilers that take in "code" in human-readable text format and output a list of instructions. The program is designed in such a way that the produced list of valid instructions always carries out the task described in "code" faithfully.

The Java compiler is such a program for Java code. When a Java program is compiled, the compiler reviews the code for syntax errors and then translates it into bytecode, a type of intermediate language closer to machine language. The Java Virtual Machine (JVM) then interprets this bytecode into machine code that your computer's processor can execute.

Syntax errors or other kinds of errors essentially refer to situations where the compiler doesn't know how to, or cannot produce a valid set of instructions that can carry out what the code is describing.

Compilers like the Java compiler don't just translate code. They also optimize it, making it more efficient so that the resulting program runs faster and consumes less memory. For example, they may look at your `for` loop that is adding up the first N natural numbers and decide to replace the loop with the formula for this computation instead.

2.5 Operating Systems

With a sound understanding of the fundamental components of modern computing, it's important to highlight the role of the operating system.

An operating system (OS) is a type of system software that manages computer hardware and software resources and provides various services for computer programs. It acts as a mediator between users and the computer hardware. Users interact with the operating system through user interfaces such as a command-line interface (CLI) or a graphical user interface (GUI).

Operating systems bear the responsibility of managing the computer's resources, including the processor, memory, disk space, and input/output devices. They coordinate tasks, ensuring that the processor's time is used judiciously, and manage the memory, keeping track of which parts are in use and which are available.

In essence, the operating system provides the platform on which all other software runs. It is the environment in which programs, written in languages like Java and then compiled, operate.

This fundamental understanding of modern computing components helps elucidate the intricate operations that are continuously happening within our laptops, desktops, and even our smartphones. These fundamental aspects form the backbone of the digital age we live in.

Part II

Algorithmic Analysis

3 Algorithmic Analysis

In this chapter, we will delve into the exciting world of algorithmic analysis. The objectives of this chapter are to learn to communicate the speed of an algorithm effectively and to appraise the performance of an algorithm from pseudocode or Java code.

3.1 Problems, Algorithms, and Programs

Before understanding algorithmic analysis, it's essential to differentiate between problems, algorithms, and computer programs. These are three distinct concepts that are interrelated.

3.1.1 Problems

A problem in computer science refers to a specific task that needs to be solved. It can be thought of in terms of inputs and matching outputs. For instance, to solve the problem of finding the youngest student in our class, the input would be the names and ages of all students in the class. The output would be the name of the youngest student.

It's helpful to perceive problems as functions in a mathematical sense. In mathematics, a function is a relationship or correspondence between two sets — the input set (domain) and the output set (range).

3.1.2 Algorithms

An algorithm, on the other hand, is a method or a process followed to solve a problem. If we perceive the problem as a function, then an algorithm can be seen as an implementation of this function that transforms an input into the corresponding output.

Since there are typically numerous ways to solve a problem, there could be many different algorithms for the same problem. Having multiple solutions is advantageous because a specific solution might be more efficient than others for certain variations of the problem or specific types of inputs.

For instance, one sorting algorithm might be best suited for sorting a small collection of integers. Another might excel in sorting a large collection of integers, while a third might be ideal for sorting a collection of variable-length strings.

By definition, a sequence of steps can only be called an algorithm if it fulfills the following properties:

- It must be correct.
- It consists of a series of concrete steps.
- There is no ambiguity about the step to be performed next.
- It must comprise a finite number of steps.
- It must terminate.

3.1.3 Programs and Their Building Blocks

Before discussing programs in detail, let's briefly review two essential components that make a program run: the CPU (Central Processing Unit) and memory.

A CPU is the electronic circuitry within a computer that has the ability to execute certain instructions. Its primary function is to fetch, decode, and execute instructions. It has slots to store data, referred to as registers. Communicating with a CPU involves telling it what operation you want to perform and on which data. For instance, you might instruct the CPU to add two numbers stored in registers A and B and store the result in register C.

CPU instructions are binary codes that specify which operation the CPU should perform. Here's an example of what they look like:

```
1001001100110011110000111011111
```

Some bits in the instruction form the **opcode**, the operation code. The opcode is a unique identifier for an operation, like adding integers. Other bits form the **operand(s)**, the data on which to operate. The operand can be where the data is stored (the name of a register or an address in memory), or where to store the result of the operation (again, the name of a register or an address in memory).

For human readability, there are notations to represent these binary instructions. Here is an example of a set of instructions in a human-readable form:

```
.global main
main:
    addi    sp, sp, -16
    sd     t0, 0(sp)
    sd     t1, 8(sp)
    call   some_function
    ld     t0, 0(sp)
    ld     t1, 8(sp)
# Use t0 and t1 here as if nothing happened.
```

```
addi    sp, sp, 16
```

Programs are structured into sections. They include code sections, which contain a list of instructions, and data sections, which hold binary data such as text, images, or numbers that the program needs to use. Typically, there is a designated “main” section that contains the instructions to be executed first.

When you initiate an executable (with the exception of Mac “applications”), the binary data (“bits”) are read from the hard disk and transferred to the main memory (RAM). The execution of the program begins when the first instruction from the “main” section is transferred to the CPU.

In this context, a computer program’s code can be seen as an instance, or concrete representation, of an algorithm. Although the terms “algorithm” and “program” are distinct, they are often used interchangeably for simplification.

3.2 Comparing the Performance of Programs

When you compile or build a program, its code is converted into a series of instructions and data in memory. However, the execution time of the same program can vary across different machines due to differences in the processor’s capabilities.

As each machine can potentially have a different processor, comparing the speed of programs can be a complex task, and is only meaningful when the processor is the same or standardized. Even when the processor is standardized, many factors affect performance -

- Background tasks on one machine can interfere with performance measurements.
- Even if you have the exact same processor, differences in manufacturing mean each can run at a different clock frequency.
- Small differences in ambient temperature affect how high a processor can clock.
- The mounting pressure of a cooler can affect heat transfer and in turn how high a processor can clock.
- Many cooling solutions involve vapor chambers. The orientation of vapor chambers can affect heat transfer and in turn how high a processor can clock.
- Even cosmic radiation can affect processors and memory.

Therefore, we usually prefer to compare algorithms instead. The methods of comparing algorithms will be discussed in the following sections of this chapter.

3.3 Analyzing Algorithms

One of the key components of this course is to provide a framework for predicting the performance of algorithms just by inspecting their structure. Let's dive into the process of analyzing an algorithm's time complexity.

3.3.1 Predicting Execution Time

Consider a simple method that adds two numbers:

```
public int add(int lhs, int rhs)
```

Suppose calling `add(2, 4)` takes 1 second. How long would `add(40, 50)` or `add(343245634, 32432423)` take? As you might expect, all these operations, despite the difference in magnitude of the numbers involved, take approximately the same time. That's because the time complexity of an addition operation does not depend on the values of the numbers but on the number of operations involved, which, in this case, is a single addition.

3.3.2 Impact of Input Size

Now, let's examine a slightly more complex method that sums up a list of numbers:

```
public int sumOfList(List<int> l)
```

Assuming that adding two numbers takes one second, how long would summing a list of 10 numbers take? We can infer that the time taken by `sumOfList` depends on the size of the list we provide. For instance, summing up a list of 10 numbers would take about half the time needed to sum up a list of 20 numbers.

Here's an implementation of `sumOfList`:

```
import java.util.ArrayList;

class Square {
    static int sumOfList(ArrayList<Integer> l) {
        int sum = 0;
        for (int i : l) {
            sum += i;
        }
        return sum;
    }
}
```

}

The key insight is that the execution time of this method depends on the number of elements in the list - which is the size of the input. The time complexity is directly proportional to the number of times the addition statement is executed, which is equal to the size of the list.

3.3.3 Iterations and Input Size

Let's take it a step further. If you're summing a list of N items, each of which is another list of M items, the operation would take $N * M$ addition statements. Here, the time complexity is a function of both N and M .

3.3.4 Gauging Relative Execution Time

The crux of analyzing an algorithm's performance lies in understanding how many times statements in the program run as a function of the size of the input. This approach enables us to estimate how long two invocations of the same method will take relative to each other, given the size of the input for each.

Understanding this concept will allow you to better predict the performance of algorithms, which is a crucial skill in efficient programming and system design.

3.4 Algorithm Complexities

In order to evaluate an algorithm's efficiency, we analyze its time complexity and space complexity, both of which describe how the algorithm's performance scales with the size of the input.

3.4.1 Time Complexity

Time complexity measures how the execution time of an algorithm increases with the size of the input. For instance, counting how many times the number 5 appears in a list requires checking each number in the list. Hence, the time complexity is directly proportional to the size of the list.

3.4.2 Space Complexity

Space complexity quantifies the amount of memory an algorithm requires relative to the size of the input. Using the previous example, we would need enough space to store the list and an additional space to store the counter. Thus, the space complexity is proportional to the size of the list, or more precisely, $N + 1$.

3.4.3 Conditional Statements and Complexity

In cases where conditional statements are present, the number of executed statements depends on which branch the program takes. One branch may contain more statements than the other. To handle such scenarios, we introduce the concept of Big-O, Big- Ω , and Big- Θ notations.

3.4.3.1 Big-O Notation

The Big-O notation describes the worst-case time complexity of an algorithm, essentially providing an upper bound on the time taken. This notation considers the scenario where the program consistently takes the path with the most statements. While Big-O is commonly used for time complexity, it can also describe space complexity.

The notation comprises two parts: the function itself and the variable representing the input size. Generally, ‘n’ is used to represent the input size, and constants and coefficients are typically ignored. For instance, the following functions are all $O(n)$:

- $n + 1$
- $2n$
- $103n + 124$

Only the highest degree of ‘n’ is considered when determining Big-O notation. Therefore, functions such as n^2 , $n/2$, or \sqrt{n} are not considered $O(n)$.

3.4.3.2 Big- Ω Notation

The Big- Ω notation represents the best-case complexity of an algorithm. It follows the same format as Big-O notation but focuses on the scenario where the program consistently takes the path with the fewest statements.

3.4.3.3 Big-Θ Notation

The Big-Θ notation is used to denote the average-case complexity of an algorithm. It again follows the same structure as Big-O, but it considers both the best and worst-case scenarios to provide an average estimate of the algorithm's performance.

Remember, these notations and complexities are pivotal in estimating the performance of an algorithm based on the size or other properties of the input, correlating to the number of steps taken by the algorithm. This understanding is crucial when designing efficient and effective solutions in computer science.

3.5 Common Big-O Complexities

Big-O notation is a way of expressing the worst-case time complexity of an algorithm. It describes how the running time of an algorithm changes as the size of its input grows. The most common Big-O complexities are:

3.5.1 Constant Complexity and Growth

Constant complexity, often represented as $O(1)$, occurs when the running time of an algorithm or the amount of work needed does not change with the size of the input (N). This means the algorithm takes a fixed amount of time, regardless of how many elements it is processing.

For example, accessing an element in an array by its index is an operation of constant complexity. This is because it takes roughly the same amount of time, regardless of the size of the array:

- $f(n) = 1$ for all n

In this case, no matter how large or small our input size is, the amount of work we have to do remains the same. This is the most efficient complexity an algorithm can have.

3.5.2 Linear Complexity and Growth

Linear complexity, often represented as $O(n)$, occurs when the running time of an algorithm or the amount of work needed scales proportionally with the size of the input (N). For every additional element in the input, a fixed amount of work is added.

For example, finding an element in an unsorted list is an operation of linear complexity, as the algorithm might need to look at every element once:

- $f(n) = n$

In this case, if we add one more element to our input size, we add one more unit of work. This is because every additional element requires the same amount of work.

3.5.3 Quadratic Complexity and Growth

Quadratic complexity, often represented as $O(n^2)$, occurs when the running time of an algorithm or the amount of work needed scales with the square of the size of the input (N). For every additional element in the input, the work increases by a factor of n .

For example, a simple nested loop for comparing pairs of elements in a list has quadratic complexity. This is because each element is compared to every other element:

- $f(n) = n^2$

In this case, if we add one more element to our input size, we add n units of work, as we have to compare this new element with every other element already in the list.

3.5.4 Exponential Functions and Growth

An exponential function is one that includes a variable in the exponent. To illustrate, the function 2^n is an exponential function. Here, with each unit increase in the input, the output is multiplied (or scaled up) by a factor of 2. For instance:

- $f(1) = 2^1 = 2$
- $f(2) = 2^2 = 4$
- $f(n + 1) = f(n) * 2$

Exponential growth is characterized by a constant factor scaling up the running time of the algorithm, or the amount of work needed, for each unit increase in the size of the input (often denoted as N). This constant factor is the base of the exponent. This means that for an algorithm with a time complexity of 2^N , adding one more element to the input could potentially double the amount of work required. Examples of algorithms exhibiting exponential growth include certain solutions to problems like the Towers of Hanoi and calculations of the Fibonacci sequence.

3.5.5 Logarithmic Functions and Growth

Logarithmic functions serve as the inverse of exponential functions. For example, $\log_2(n)$ is a logarithmic function. Here, the output is decremented by 1 each time the input is divided (or scaled down) by a factor of 2. For instance:

- $f(16) = \log_2(16) = 4$

- $f(8) = \log_2(8) = 3$
- $f(n/2) = f(n) - 1$

Logarithmic growth is characterized by a decrement of 1 in the running time of the algorithm or the amount of work needed, each time the size of the input (denoted as N) is divided by a constant factor. This constant factor is the base of the logarithm. For an algorithm with a time complexity of $\log_2(n)$, if we have an input size of 16, doubling the input size will only increase the amount of work by 1 unit. Similarly, halving the input size will decrease the work by 1 unit.

Algorithms often display complexities such as $N * \log_2(N)$, which represents a combination of linear and logarithmic growth. Examples of such algorithms that include logarithmic growth are binary search and certain sorting algorithms.

3.5.6 Examples

Let's look at some examples of algorithms and their Big-O complexities.

```
static int findMin(x, y) {
    if (x < y) {
        return x;
    } else {
        return y;
    }
}
```

This algorithm finds the minimum of two numbers x and y . It does not depend on the input size, since it only performs one comparison and one return statement. Therefore, its worst-case complexity is $O(1)$. Its best case and average case are also $O(1)$ since they are the same as the worst case.

```
static int linearSearch(numbers[], target)
    for (int i = 0; i < numbers.length; i++) {
        if (numbers[i] == target) {
            return i;
        }
    }
    return -1;
}
```

This algorithm performs a linear search on an array of numbers to find a target value. It iterates through each element of the array until it finds the target or reaches the end of the array. In the worst case, it has to check every element of the array, which means its worst-case

complexity is $O(n)$, where n is the length of the array. In the best case, it finds the target in the first element, which means its best-case complexity is $O(1)$. In the average case, it finds the target somewhere in the middle of the array, which means its average-case complexity is also $O(n)$.

We will talk about more examples of other common worst-case complexities throughout this course.

3.6 Growth Rate

The growth rate for an algorithm is the rate at which the cost of the algorithm grows as the size of its input grows. The cost can be measured in terms of time, space, or other resources. The worst-case complexity notation essentially denotes the growth rate of the time complexity with respect to the size of a worst-case input.

The table below summarizes how different Big-O complexities compare in terms of their growth rates.

Complexity	Growth Rate
$O(1)$	Constant
$O(\log(n))$	Logarithmic
$O(n)$	Linear
$O(n^2)$	Quadratic
$O(2^n)$	Exponential

As we can see, constant and logarithmic complexities have very low growth rates, meaning that they are very efficient and scalable algorithms. Linear complexity has a moderate growth rate, meaning that it can handle reasonably large inputs but may become slow for very large inputs. Quadratic and exponential complexities have very high growth rates, meaning that they are very inefficient and unscalable algorithms that can only handle small inputs.

3.7 Summary / Review

In this section, we learned how to analyze the performance of algorithms using Big-O notation. We saw that measuring the actual running time of an algorithm on real hardware is difficult and impractical because it depends on many factors such as the hardware specifications, the programming language, the compiler, the input size, and the input distribution. Therefore, we need a way to simplify and standardize the performance analysis across different hardware and software platforms.

We learned that one way to simplify the performance analysis is to count the number of steps or instructions that an algorithm needs to execute before finishing. We saw that these steps are analogous to the instructions generated by a compiler when it translates our code into machine code. We also learned that different steps may have different costs depending on their complexity, but we can ignore these differences for simplicity and focus on the overall number of steps.

We learned that Big-O notation is a mathematical tool that allows us to express the worst-case complexity of an algorithm. It describes how the number of steps grows as a function of the input size in the worst possible scenario. It gives us an upper bound on the performance of an algorithm, meaning that it tells us how slow an algorithm can get in the worst case. We also learned that Big-O notation ignores constant factors and lower-order terms because they become insignificant as the input size grows.

We learned about some common Big-O complexities and their growth rates, such as constant, logarithmic, linear, quadratic, and exponential. We saw some examples of algorithms and their Big-O complexities, and how to analyze them using simple rules such as adding complexities for sequential steps, multiplying complexities for nested steps, and taking the maximum complexity for conditional steps.

We learned that Big-O notation helps us compare different algorithms and choose the most efficient one for a given problem. It also helps us estimate how well an algorithm can scale to larger inputs and how it can affect the performance of our applications.

Part III

Unit Testing

4 Unit Testing and Test-Driven Development

4.1 Background

Imagine you are an architect and you've just designed a large, complex building - let's say a skyscraper. This skyscraper is not merely a single entity; it's composed of thousands of individual components - the plumbing, the electrical wiring, the elevators, the heating system, and the building's structural elements, among many others. Now, how would you ensure that the skyscraper works as intended? Would you wait until the entire building is constructed and then start testing every possible scenario? Obviously, this approach is time-consuming, and it exposes you to significant risk.

A similar challenge exists in the world of software development. Take, for example, a web browser. This is a complex piece of software with hundreds of classes interacting in intricate ways. These classes and methods perform various tasks such as rendering HTML and CSS, processing JavaScript, managing cookies, implementing security features, and many others. Ensuring the correct functionality of this software is a daunting task, given the vast range of potential inputs. After all, there are billions of web pages on the internet, each with its unique combination of technologies, designs, and user interactions. How can you guarantee that your browser works flawlessly with all of them? A naive approach would be to load each web page and observe the output, but this process is not only time-consuming but also practically impossible.

This conundrum begs the question: How can we validate the correct functionality of a software product efficiently? The direction points towards automation - the ability to conduct tests without manual intervention. But how can we achieve this, especially given the enormous application surface area?

This is where unit testing and Test-Driven Development (TDD) come in.

4.2 The Power of Unit Testing

While the surface area of an entire application is vast, the surface area of individual classes and units of code within the software project is significantly smaller. If we can write tests to verify that each method within each class functions correctly for all possible inputs, we reduce the complexity of the problem.

At first glance, it might seem like an overwhelming task. Even a moderately complex software project can have thousands of methods spread across hundreds of classes. Writing tests for all of them could result in thousands of test cases. But this is precisely where automation proves its worth. By automating these tests, we can execute them each time we modify our code, ensuring the functionality remains intact. This method gives us confidence that our changes have not inadvertently introduced bugs into existing functionality.

The key principle here is that by ensuring each individual unit of our software behaves correctly, we can be reasonably confident that the application as a whole operates as expected, provided the software architecture is sound. In this manner, unit testing allows us to break down the monumental task of verifying a complex software system's functionality into manageable, automated tasks.

In the following sections, we will delve deeper into the concept of unit testing, its implementation in Java, and the practice of Test-Driven Development, where tests actually guide and shape the development of the software. Buckle up, for we're about to embark on an exciting journey that will fundamentally change how you approach software development!

4.3 A Basic Approach to Unit Testing

In order to illustrate the process of unit testing in Java, let's consider a simple utility class named `MathUtil`. This class defines basic arithmetic operations such as `add`, `subtract`, etc.

```
public class MathUtil {  
    public int add(int a, int b) {  
        return a + b;  
    }  
  
    // More methods for subtract, multiply, etc.  
}
```

As we discussed in the previous section, to ensure our `MathUtil` functions correctly, we associate it with a `MathUtilTest` class. This class contains multiple test methods, each designed to verify a different scenario of the operations provided by `MathUtil`.

```
public class MathUtilTest {  
  
    public boolean testAdd1() {  
        MathUtil m = new MathUtil();  
        int lhs = 5;  
        int rhs = 7;
```

```

        if (m.add(lhs, rhs) == lhs + rhs) {
            return true;
        } else {
            return false;
        }
    }

    // More test methods for other cases...
}

```

In the above example, the `testAdd1` method tests the addition of two positive numbers. We could also add methods like `testAdd2` to test adding a positive and a negative number, `testAdd3` to test adding two negative numbers, and so forth. Each of these methods tests a specific scenario and validates that the result is as expected.

4.4 Recognizing the Inefficiencies and Redundancies

While the above approach accomplishes our objective of validating the methods in our `MathUtil` class, you might have already noticed that it's far from optimal. There are several glaring issues that can make this method tedious and inefficient:

4.4.1 Redundancy

Every test method follows a similar pattern - we perform an operation and then verify if the result matches the expected outcome. This redundancy suggests we could abstract out the verification part into a separate method.

In the following expanded `MathUtilTest` class, you can observe that `testAdd2` and `testAdd3` follow the exact same pattern as `testAdd1`. They create an instance of `MathUtil`, perform an operation, and then compare the result with the expected outcome. This repetitive pattern across multiple tests highlights the redundancy and inefficiency of this approach.

```

public class MathUtilTest {

    public boolean testAdd1() {
        MathUtil m = new MathUtil();
        int lhs = 5;
        int rhs = 7;

        if (m.add(lhs, rhs) == lhs + rhs) {

```

```

        return true;
    } else {
        return false;
    }
}

public boolean testAdd2() {
    MathUtil m = new MathUtil();
    int lhs = -5;
    int rhs = 7;

    if (m.add(lhs, rhs) == lhs + rhs) {
        return true;
    } else {
        return false;
    }
}

public boolean testAdd3() {
    MathUtil m = new MathUtil();
    int lhs = -5;
    int rhs = -7;

    if (m.add(lhs, rhs) == lhs + rhs) {
        return true;
    } else {
        return false;
    }
}

// More test methods for other cases...
}

```

4.4.2 Lack of Automation

Let's see how we need to currently run the tests we've written.

```

public class MathUtilTest {

    // testAdd1, testAdd2, testAdd3, etc. test methods...
}

```

```

public static void main(String[] args) {
    MathUtilTest test = new MathUtilTest();

    System.out.println("testAdd1 result: " + (test.testAdd1() ? "PASS" : "FAIL"));
    System.out.println("testAdd2 result: " + (test.testAdd2() ? "PASS" : "FAIL"));
    System.out.println("testAdd3 result: " + (test.testAdd3() ? "PASS" : "FAIL"));

    // add more prints for other test cases
}
}

```

With this `main` method, you can now run the `MathUtilTest` class, and it will execute each of the `testAdd` methods and print whether each test passed or failed. This method is a basic way to manually execute the tests and check their results.

Currently, we need to call each test method manually to run our tests. An automated system that could execute all tests for us would save time and reduce the chances of human error. However, as we will see later, there are better approaches to automation than the one we've used here.

4.5 A slightly more sophisticated approach to Unit Testing

To alleviate the redundancy, we can create a method that compares the expected and actual results and raises an error if they do not match. This method, which we can call `assertEquals`, would look something like this:

```

public static void assertEquals(String testCaseName, int expected, int actual) {
    if (expected != actual) {
        System.out.println(testCaseName + " result: FAIL");
    } else {
        System.out.println(testCaseName + " result: PASS");
    }
}

```

Then we can simplify our test methods by using `assertEquals`:

```

public void testAdd() {
    MathUtil m = new MathUtil();

    assertEquals("testAddTwoPositive", m.add(5, 7), 13);
    assertEquals("testAddTwoNegative", m.add(-5, -7), -13);
}

```

```
        assertEquals("testAddNegPos", m.add(-5, 7), 2);
    }
```

Now, our test case looks cleaner and easier to understand. The `assertEquals` method abstracts away the comparison details, leaving only the test logic in the test case. We can apply this simplification to all our test methods.

This approach significantly reduces the redundancy in our test code, making it easier to write and maintain our tests. However, we are still manually running each test method from the `main` method. What if we could also automate the execution of all test methods?

4.6 Automating Test Execution

What if we could just call a single method that runs all our test methods? Let's define a simple `runTests` method that does exactly that.

```
public void runTests() {
    testAdd();
    // Call all other test methods here...
}
```

And then you can simply call the `runTests` method to execute all your tests:

```
public static void main(String[] args) {
    MathUtilTest test = new MathUtilTest();
    test.runTests();
}
```

This approach is an improvement over manually running each test. However, it still has some drawbacks. For instance, when you add a new test method, you need to remember to add a call to this method in the `runTests` method. It would be better if our test framework could automatically detect and run all test methods without requiring any modifications to the `runTests` method. As we'll see later, this is precisely what test frameworks like JUnit offer.

4.7 Summary

So far, we have seen how unit testing can be a powerful tool in ensuring that individual units of code within a larger software application function as expected. We have also discussed and

implemented a basic system for automating unit tests in Java, gradually refining this system to make it more efficient and less redundant.

In the following sections, we will discuss JUnit, a popular unit testing framework in Java that takes automation and convenience to the next level. We will also explore the practice of Test-Driven Development, where we let our tests guide the development of our software, helping us to write cleaner, more robust code. Stay tuned!

5 Writing JUnit Tests and Test-Driven Development

Unit testing is crucial to ensure the accuracy and performance of your code. But as we've seen, managing and writing tests can be a bit cumbersome. This is where testing frameworks, like JUnit, come in. They automate the tedious parts of testing and provide us with a plethora of tools to write effective tests.

5.1 Introduction to JUnit

JUnit is a widely used testing framework in the Java world. It automates the process of running tests and provides us with a wide range of assertion methods to validate our code. JUnit helps to simplify our test code, making it easier to read and maintain.

So, why is JUnit so popular?

1. **Simplicity:** JUnit simplifies the process of writing and running tests. The framework handles the boilerplate code, allowing us to focus solely on writing the test cases.
2. **Assertion Library:** JUnit provides a comprehensive set of assertion methods that help us validate our code against a wide range of conditions.
3. **Annotations:** JUnit uses annotations to define test methods and setup methods, making our test code easier to read and understand.
4. **Automatic Test Discovery:** JUnit automatically finds and runs all test methods, so we don't have to manually list them in our code.
5. **IDE Integration:** Most modern IDEs provide first-class support for JUnit, including features such as generating test cases and displaying test results in a friendly format.

Now that you understand what JUnit is and why it's beneficial let's see how to use it in our `MathUtil` class.

5.2 Useful JUnit Assertions

JUnit provides a set of methods called assertions that are used to test the expected output of your code. These assertions help verify that your code behaves as expected under different conditions.

Let's take a look at some commonly used assertions:

5.2.1 assertEquals

This assertion checks if two values are equal:

```
assertEquals(expected, actual);
```

If `actual` is not equal to `expected`, the assertion fails, and the test method will terminate immediately.

Let's rewrite our `addTest1` method using JUnit's `assertEquals`:

```
@Test
public void addTest1() {
    MathUtil m = new MathUtil();
    int lhs = 5;
    int rhs = 7;

    assertEquals(lhs + rhs, m.add(lhs, rhs));
}
```

5.2.2 assertTrue and assertFalse

These assertions verify if a condition is `true` or `false`, respectively:

```
assertTrue(condition);
assertFalse(condition);
```

If the `condition` does not meet the expectation (i.e., `true` for `assertTrue` and `false` for `assertFalse`), the assertion fails, and the test method will terminate immediately.

5.2.3 `assertNotNull` and `assertNull`

These assertions check if an object is `null` or not:

```
assertNotNull(object);
assertNull(object);
```

If the `object` does not meet the expectation (i.e., not `null` for `assertNotNull` and `null` for `assertNull`), the assertion fails, and the test method will terminate immediately.

These are just a few examples. JUnit provides a comprehensive set of assertions to cover almost any condition you might want to verify.

5.3 Setting Up JUnit

Before you can use JUnit, you need to make sure the library is on your classpath. This process can vary depending on the IDE and build system you're using.

For our labs, we will ensure the `JUnit` library is on our classpath by pre-configuring the project and IDE for you. However, if you're working on your own project, you'll need to add the `JUnit` library to your project's classpath.

When working on your own projects, you might be interested in using a build system like [Maven](#) or [Gradle](#) to manage your dependencies. These build systems make it easy to add and manage dependencies in your project. For example, if you're using Maven, you can add the following dependency to your `pom.xml` file:

```
<dependency>
    <groupId>org.junit.jupiter</groupId>
    <artifactId>junit-jupiter-api</artifactId>

    <version>5.7.0</version>
    <scope>test</scope>
</dependency>
```

This will automatically download the JUnit library and add it to your project's classpath.

Regardless of how you added the `JUnit` library to your project, next, we need to import the necessary classes and annotations from JUnit. At the top of our `MathUtilTest` class, we add the following import statements:

```
import static org.junit.jupiter.api.Assertions.*;
import org.junit.jupiter.api.Test;
import org.junit.jupiter.api.BeforeEach;
```

The first import statement statically imports all assertion methods from `Assertions`, allowing us to use them directly in our code. The second import statement imports the `Test` annotation, which we use to denote our test methods. The third import statement imports the `BeforeEach` annotation, which we'll discuss in a moment.

5.4 Utilizing `@Test` and `@BeforeEach` Annotations

In JUnit, we use the `@Test` annotation to indicate that a method is a test method. This allows JUnit to automatically discover and run this method as a test.

```
@Test
public void addTest1() {
    // test code...
}
```

However, what if we have some setup code that we want to run before each test? This is where the `@BeforeEach` annotation comes in. Any method annotated with `@BeforeEach` will be run before each `@Test` method.

Let's say we want to create a new `MathUtil` instance before each test:

```
MathUtil m;

@BeforeEach
public void setup() {
    m = new MathUtil();
}
```

Now, before each test method is run, JUnit will first execute the `setup` method, ensuring that we have a fresh `MathUtil` instance for each test.

5.5 Interpreting JUnit Test Runner Output

Understanding the output of the JUnit test runner is crucial for interpreting the results of your tests. This helps you diagnose issues in your code and identify exactly what went wrong.

Let's analyze the output of the `junit-platform-console-standalone` test runner to get a feel for how this works.

```
JUnit Jupiter
JUnit Vintage
    BinarySearchTreeHiddenTest
        testInsertAndSearch
        testDeleteSingleNode
        testTreeTraversals  expected:<[2, 3, 4, 5, [6, 7], 8]> but was:<[2, 3, 4, 5, [7, 6]
        testContainsElementNotInTree
        testContainsEmptyTree
        testDeleteEmptyTree
        testContainsElementInTree
        testSearchElementInTree
        testInsertMultipleElements
        testDeleteDuplicateElements
        testDeleteElementNotInTree
        testInsertNegativeNumbers
        testInsertAndSize
        testSearchEmptyTree
        testDeleteNodeWithMultipleElements
        testInsertDuplicatesAndRemove
        testInsertSingleElement
        testSearchElementNotInTree
    JUnit Platform Suite

Failures (1):
JUnit Vintage:BinarySearchTreeHiddenTest:testTreeTraversals
=> org.junit.ComparisonFailure: expected:<[2, 3, 4, 5, [6, 7], 8]> but was:<[2, 3, 4,
   DataStructures.BinarySearchTreeHiddenTest.testTreeTraversals(BinarySearchTreeHidden
   [...]
```

The JUnit console output provides a tree structure representing the test execution. The topmost nodes represent the test engines used, in this case, `JUnit Jupiter` and `JUnit Vintage`. Underneath each engine are the individual test classes, such as `BinarySearchTreeHiddenTest`.

Within each test class node, there are child nodes representing each test method, such as `testInsertAndSearch` or `testDeleteSingleNode`. These methods are marked with a `+` symbol if they passed, and with a `-` symbol if they failed. In this case, we see that `testTreeTraversals` has failed.

Accompanying the failure symbol is a brief description of the failure, which is the assertion

message from the test method. In this example, the test expected the array [2, 3, 4, 5, 6, 7, 8], but received the array [2, 3, 4, 5, 7, 6, 8]. This discrepancy caused the test to fail.

After the tree structure, there is a section titled **Failures** which provides more detailed information about each failure. For each failure, it lists:

1. The test class and method that failed.
2. The type of assertion failure that occurred, which is `org.junit.ComparisonFailure` in this case.
3. The detailed assertion failure message, which is the same as what's shown in the tree structure.
4. The location in the code where the failure occurred, which can be very useful (and is often the first thing you should look at when debugging a test failure). In this case, the failure occurred on line 220 of `BinarySearchTreeHiddenTest.java` (see the `DataStructures.BinarySearchTreeHiddenTest.testTreeTraversals(BinarySearchTreeHiddenTest)` line).

5.6 Conclusion

In conclusion, JUnit simplifies the process of writing and managing tests. It provides a comprehensive set of assertion methods to verify our code and uses annotations to define and organize our tests, making them easier to read and understand. By taking advantage of these features, we can write more effective and maintainable tests. In the next section, we'll dive deeper into Test-Driven Development, a methodology that leverages the power of testing to guide and improve the development process.

5.7 Test-Driven Development

Test-Driven Development (TDD) is a software development methodology that is centered around the idea of writing tests before writing the actual code. It is a highly disciplined process that follows a strict order of operations: red, green, refactor. This method has profound implications on the design, quality, and reliability of the software.

Let's dive into what these steps entail.

1. **Red:** Write a test that covers a specific functionality you want to implement. This test should fail initially because you haven't written the actual code yet. This stage helps you think about the functionality in detail, ponder on the inputs and expected outputs, and outline the structure of your code.

2. **Green:** Write the minimal amount of code needed to pass the test. At this stage, don't worry about the elegance of your code. Your primary focus is on functionality. Run your test, and it should pass this time.
3. **Refactor:** Refactor the code you just wrote in the green stage to eliminate duplication, improve readability, and ensure the code adheres to the best practices. After refactoring, all tests should still pass. If a test fails, it means the refactoring broke the functionality, and you need to revise your changes.

This cycle repeats for every small chunk of functionality you add to your software. With this approach, you are incrementally building your software with the assurance that at each step, the implemented functionality is working as expected.

5.7.1 The Motivation Behind Test-Driven Development

You might be wondering, why would you want to put in the extra effort to write tests before writing the actual code? Here are a few motivating factors:

1. **Confidence:** With TDD, you can be confident that your code works because you have tests that prove it. This confidence is especially important when you need to modify your code later. Changes can break existing functionality, but with a robust set of tests, you can quickly catch and fix these regressions.
2. **Better Design:** Writing tests first forces you to think about your code from a user's perspective. This shift in viewpoint often results in better code organization and modularity because you design your code to be easy to test, which typically means it is also easy to use and modify.
3. **Documentation:** Tests act as a form of documentation that shows how the code is supposed to work. New team members can look at the tests to understand what each function is supposed to do and what edge cases it handles.
4. **Development Speed:** While TDD might seem to slow you down at the beginning, it typically results in faster development in the long run. With TDD, you spend less time debugging and fixing bugs because you catch them early in the development process, before they become entangled with other parts of the code.

In conclusion, TDD is a powerful methodology that can significantly improve the quality of your code and your efficiency as a developer. While it might seem difficult at first, with practice, it becomes a natural part of the development process.

6 Introduction to Mutation Testing

6.1 Unit Testing, a Recap

Writing *correct* code is challenging. Owning your code, maintaining it, and ensuring it always performs as expected can be even more difficult. After all, we are human, and mistakes are a part of our nature. Therefore, we can't always rely solely on our ability to catch these mistakes. Instead, automated processes that can help us identify and correct these errors are truly invaluable.

Unit testing is one such automated process that tests your code. These tests target individual units of your code and are designed to verify their behavior. In programming, a method can be thought of as the smallest unit of code, and Classes, Packages, and Modules are larger units of code. Every unit test is designed to validate a single behavior of a single unit of code.

Unit testing is so crucial that they are a mandatory requirement for many software development projects. For instance, if you are contributing a Java source class to most projects, you are also required to provide unit tests for that class. In this course, where we deal with Data Structures and Algorithms, learning how to implement them in Java, etc., submitting your implementations without Unit tests would be like submitting a paper without a bibliography. It might make us question, “Sure, you wrote this and it looks convincing, but where’s your evidence? Why should I believe you?”

And that's why you will be required to write unit tests for your Data Structures and Algorithms implementations. But then, how do you know that your unit tests are correct?

Consider this example:

```
@Test  
public void testAdd() {  
    assertTrue(1 == 1);  
}
```

The test passes, but does it mean it's correct? Let's examine the test.

6.2 Writing *Correct* Unit Tests

Taking a closer look at the unit test above, two problems become evident:

- The `add` method isn't called in the test.
- The test can't fail.

To address the first issue, we can leverage tools that detect when a unit test doesn't cover a particular line of code. Here "cover" means "execute".

So what about a test like this?

```
@Test  
public void testAdd() {  
    assertTrue(add(1, 1) == add(1, 1));  
}
```

The `add` method is called, and test coverage is at 100%. However, this test can't fail, meaning it is not *correct*.

6.3 Mutation Testing: A Solution to Test the Correctness of Your Unit Tests

Mutation testing provides a way to test the correctness of your unit tests. It can help you find tests that can't fail, tests that need more test cases, and even logic errors in your code.

The basic idea of mutation testing is as follows: if you claim your source code is correct, and that the tests prove it, mutation testing will challenge that claim. Mutation testing introduces small changes, called *mutations*, to your source code. If your tests still pass after these changes, it suggests that your test or source code may not be correct.

These mutations are created by algorithms called *mutators* or *mutation operators*. Each time a mutator runs, it receives a fresh copy of your source code and makes only one change. The result of a mutator is a *mutant* - a mutated version of your source code.

6.4 Mutation Testing in Practice

Here's how mutation testing works in practice:

- Apply a set of mutators to your source code to produce a collection of mutants.
- For each mutant, run your unit tests.

- If the unit tests pass, the mutant has *survived*.
- If the unit tests fail, the mutant was *killed*

The aim is to kill as many mutants as possible. If too many mutants survive, it indicates that your unit tests are not sufficient to prove that your source code is correct. If you manage to kill all mutants, you can use your unit tests to argue that your source code is indeed correct.

In our course, only unit tests that leave no surviving mutants will be accepted as a valid submission.

6.4.1 What are these mutators?

Mutators introduce specific types of changes to your code. For instance, a primitive returns mutator (PRIMITIVE_RETURNS) replaces int, short, long, char, float, and double return values with 0. So, for example, this method:

```
public int add(int lhs, int rhs) {
    return lhs + rhs;
}
```

becomes:

```
public int add(int lhs, int rhs) {
    return 0;
}
```

This mutant would survive in two cases: if the `add` method is never tested or if the only test cases for `add` are ones where the result is 0. To fix this, you can add more test cases that test `add` with non-zero results.

Another example of a mutator is the Remove Conditionals Mutator (REMOVE_CONDITIONALS), which removes all conditional statements such that the guarded statements always execute.

For example, this code:

```
if (a == b) {
    // do something
}
```

becomes:

```
if (true) {  
    // do something  
}
```

This mutant would survive if the source method is never tested, if all test cases for the source method only ever test the true case, or if both branches have the exact same code or equivalent code, to begin with. To fix this, you can add more test cases that test the false case, and ensure that the true and false branches are not equivalent.

Through mutation testing, you will get valuable feedback to improve your code and tests, making them more robust and reliable.

Remember, a well-tested code base is not just about coverage—it’s about the quality of tests, their ability to catch mistakes, and their resilience against possible errors. Mutation testing is an invaluable tool in achieving this.

Certainly, let’s decipher the feedback from the Coding Rooms.

6.5 Understanding CodingRooms Feedback

In the feedback provided by the autograder in Coding Rooms, you can find two main parts. The first part lists the mutation tests that have been run, and the second part provides a grading overview.

6.5.1 Understanding Mutation Test Feedback

Each mutation test feedback starts with “Running Mutation tests”, followed by details about each mutation test performed.

```
Running Mutation tests -  
Ran mutation tests for Calculator.CalculatorTest -  
-[ RECORD 0 ]-----+-----  
Mutation type      | RemoveConditionalMutator_EQUAL_ELSE  
Source method mutated | divide  
Line no. of mutation | 56  
Test examined       | None  
Result             | SURVIVED
```

In the example above, the autograder has run a mutation test on the “divide” method of your “Calculator” class. Here is what each line in the record means:

- **Mutation type:** The type of mutation that was made to your code. In this case, a RemoveConditionalMutator_EQUAL_ELSE mutation was made. This mutation type removes a conditional (==) operator and always takes the else branch.
- **Source method mutated:** The method in your code that was mutated for the test. In this case, it's the “divide” method.
- **Line no. of mutation:** The line number in your code where the mutation was applied.
- **Test examined:** This line indicates which test case was run against the mutant. “None” means no specific test case was chosen, and all available tests were run.
- **Result:** The outcome of the mutation test. If your tests fail against the mutated code (which is a good thing!), this line will read “KILLED”. If your tests pass (which implies your tests didn't catch the error introduced by the mutation), this line will read “SURVIVED”. In this case, the mutant has survived, indicating your tests didn't catch the error.

If any mutants survive, the autograder lists those under the line, “Problematic mutation test failures printed about.”

6.5.2 Understanding the Grading Overview

The grading overview gives you a quick summary of how your code has performed in the tests. It breaks down the grading by requirements.

Grading Overview		
Requirement	Grade	Reason
1	10.00/10.00	- 4/4 tests passing.
2	32.00/40.00	-8 Penalty due to surviving mutations
Total: 42.00/50.00		

In the example above, we have:

- **Requirement 1:** This section corresponds to the first requirement of the assignment. The student received a full score (10.00/10.00) because all 4 test cases associated with this requirement passed.

- **Requirement 2:** This section corresponds to the second requirement of the assignment. The student lost 8 points because of surviving mutations, resulting in a score of 32.00/40.00 for this requirement.

The total grade of the assignment is

42.00/50.00, indicating the need for further improvement.

Remember, each surviving mutation indicates a flaw in your test cases—they didn't catch an erroneous mutation. To improve your grade, aim to update your test cases to ensure that they effectively detect the mutations.

7 List of Mutators

7.1 Conditionals Boundary Mutator (CONDITIONALS_BOUNDARY)

The conditionals boundary mutator replaces the relational operators <, <=, >, >= with their boundary counterpart as per the table below.

Original conditional	Mutated conditional
<	<=
<=	<
>	>=
>=	>

{:.table }

For example

```
if (a < b) {  
    // do something  
}
```

will be mutated to

```
if (a <= b) {  
    // do something  
}
```

7.2 Increments Mutator (INCREMENTS)

The increments mutator will mutate increments, decrements and assignment increments and decrements of local variables (stack variables). It will replace increments with decrements and vice versa.

For example

```
public int method(int i) {  
    i++;  
    return i;  
}
```

will be mutated to

```
public int method(int i) {  
    i--;  
    return i;  
}
```

Please note that the increments mutator will be applied to increments of **local variables only**. Increments and decrements of member variables will be covered by the [Math Mutator](#).

7.3 Invert Negatives Mutator (INVERT_NEGS)

The invert negatives mutator inverts negation of integer and floating point numbers. For example

```
public float negate(final float i) {  
    return -i;  
}
```

will be mutated to

```
public float negate(final float i) {  
    return i;  
}
```

7.4 Math Mutator (MATH)

The math mutator replaces binary arithmetic operations for either integer or floating-point arithmetic with another operation. The replacements will be selected according to the table below.

Original conditional	Mutated conditional
+	-
-	+
*	/
/	*
%	*
&	
	&
^	&
«	»
»	«
»>	«

{:.table}

For example

```
int a = b + c;
```

will be mutated to

```
int a = b - c;
```

Keep in mind that the + operator on **Strings** as in

```
String a = "foo" + "bar";
```

is **not a mathematical operator** but a string concatenation and will be replaced by the compiler with something like

```
String a = new StringBuilder("foo").append("bar").toString();
```

Please note that the compiler will also use binary arithmetic operations for increments, decrements and assignment increments and decrements of non-local variables (member variables) although a special **iinc** opcode for increments exists. This special opcode is restricted to local variables (also called stack variables) and cannot be used for member variables. That means the math mutator will also mutate

```
public class A {
    private int i;
```

```
public void foo() {
    this.i++;
}
}
```

to

```
public class A {
    private int i;

    public void foo() {
        this.i = this.i - 1;
    }
}
```

See the [Increments Mutator](#) for details.

7.5 Negate Conditionals Mutator (NEGATE_CONDITIONALS)

The negate conditionals mutator will mutate all conditionals found according to the replacement table below.

Original conditional	Mutated conditional
==	!=
!=	==
<=	>
>=	<
<	>=
>	<=

{:.table }

For example

```
if (a == b) {
    // do something
}
```

will be mutated to

```

if (a != b) {
    // do something
}

```

This mutator overlaps to a degree with the conditionals boundary mutator, but is less **stable** i.e these mutations are generally easier for a test suite to detect.

7.6 Return Values Mutator (RETURN_VALS)

This mutator has been superseded by the new returns mutator set. See [Empty returns](#), [False returns](#), [True returns](#), [Null returns](#) and [Primitive returns](#).

The return values mutator mutates the return values of method calls. Depending on the return type of the method another mutation is used.⁴

Return Type	Mutation
<code>boolean</code>	replace the unmutated return value <code>true</code> with <code>false</code> and replace the unmutated return value <code>false</code> with <code>true</code>
<code>int byte short</code>	if the unmutated return value is 0 return 1, otherwise mutate to return value 0

Return Type	Mutation
<code>long</code>	replace the unmutated return value <code>x</code> with the result of <code>x+1</code>
<code>float double</code>	replace the unmutated return value <code>x</code> with the result of <code>-(x+1.0)</code> if <code>x</code> is not NAN and replace NAN with 0
<code>Object</code>	replace non-null return values with <code>null</code> and throw a <code>java.lang.RuntimeException</code> if the unmutated method would return <code>null</code>

{:.table}

For example

```
public Object foo() {
    return new Object();
}
```

will be mutated to

```
public Object foo() {  
    new Object();  
    return null;  
}
```

Please note that constructor calls are **not considered void method calls**. See the [Constructor Call Mutator](#) for mutations of constructors or the [Non Void Method Call Mutator](#) for mutations of non void methods.

7.7 Void Method Call Mutator (VOID_METHOD_CALLS)

The void method call mutator removes method calls to void methods. For example

```
public void someVoidMethod(int i) {  
    // does something  
}  
  
public int foo() {  
    int i = 5;  
    someVoidMethod(i);  
    return i;  
}
```

will be mutated to

```
public void someVoidMethod(int i) {  
    // does something  
}  
  
public int foo() {  
    int i = 5;  
    return i;  
}
```

7.8 Empty returns Mutator (EMPTY RETURNS)

Replaces return values with an ‘empty’ value for that type as follows

- `java.lang.String` -> “”
- `java.util.Optional` -> `Optional.empty()`
- `java.util.List` -> `Collections.emptyList()`
- `java.util.Collection` -> `Collections.emptyList()`
- `java.util.Set` -> `Collections.emptySet()`
- `java.lang.Integer` -> 0
- `java.lang.Short` -> 0
- `java.lang.Long` -> 0
- `java.lang.Character` -> 0
- `java.lang.Float` -> 0
- `java.lang.Double` -> 0

Pitest will filter out equivalent mutations to methods that are already hard coded to return the empty value.

7.9 False returns Mutator (FALSE_RETURNS**)**

Replaces primitive and boxed boolean return values with false.

Pitest will filter out equivalent mutations to methods that are already hard coded to return false.

7.10 True returns Mutator (TRUE_RETURNS**)**

Replaces primitive and boxed boolean return values with true.

Pitest will filter out equivalent mutations to methods that are already hard coded to return true.

7.11 Null returns Mutator (NULL_RETURNS**)**

Replaces return values with null. Methods that can be mutated by the EMPTY_RETURNS mutator or that are directly annotated with NotNull will not be mutated.

Pitest will filter out equivalent mutations to methods that are already hard coded to return null.

7.12 Primitive returns Mutator (PRIMITIVE RETURNS)

Replaces int, short, long, char, float and double return values with 0.

Pitest will filter out equivalent mutations to methods that are already hard coded to return 0.

7.13 Constructor Call Mutator (CONSTRUCTOR CALLS)

Optional mutator that replaces constructor calls with `null` values. For example

```
public Object foo() {  
    Object o = new Object();  
    return o;  
}
```

will be mutated to

```
public Object foo() {  
    Object o = null;  
    return o;  
}
```

Please note that this mutation is fairly unstable and likely to cause `NullPointerExceptions` even with weak test suites.

This mutator does not affect non constructor method calls. See [Void Method Call Mutator](#) for mutations of void methods and [Non Void Method Call Mutator](#) for mutations of non void methods.

Part IV

Iterative Sorting and Searching

8 Introduction to Sorting

8.1 Understanding the Sorting Problem

Sorting is an essential problem in computer science and software engineering. Simply put, sorting is the task of rearranging a set of items, such as an array, in a specific order. For instance, an array of integers might be sorted in ascending or descending order based on their numeric value.

One might wonder why sorting is so ubiquitous in the field of software engineering. The answer lies in the multitude of applications that it has. Sorting allows for more efficient searching, data compression, and organization. It is a crucial step in numerous algorithms and computational tasks. The efficiency of a sorting algorithm can significantly affect the performance of these applications.

8.2 Developing a Sorting Algorithm from Scratch

One might find the task of developing a sorting algorithm from scratch quite daunting. Sorting algorithms, like many computational problems, have an abstract and often complex nature that can make it difficult to understand, let alone create. If shown the solution, one might understand how it works and even memorize it, but the thought process behind creating such a solution can be elusive.

A common approach to overcoming such challenges is to simplify the problem. By breaking down a complex problem into simpler subproblems, we can gain insights into the problem's structure and ultimately develop a solution. This approach will be demonstrated as we develop a sorting algorithm.

8.3 Starting Simple: Sorting an Array of Size 1

The simplest possible sorting problem involves an array of size one. But, of course, this array is already sorted as it only contains a single element. Thus, we need to do nothing to sort this array. Here's how this can be implemented in Java:

```
public static void sort1(int[] arr) {
    // do nothing
}
```

8.4 Sorting an Array of Size 2

What about an array of size two? In this case, we simply need to compare the two elements and swap them if they are out of order:

```
public static void sort2(int[] arr) {
    if (arr[0] > arr[1]) {
        int temp = arr[0];
        arr[0] = arr[1];
        arr[1] = temp;
    }
}
```

8.5 Expanding the Problem: Sorting an Array of Size 3

When sorting an array of size three, we perform similar steps as we did for the array of size two. First, we compare and potentially swap the first two numbers. Then, we do the same for the second and third numbers. Finally, we compare the first two numbers once again:

```
public static void sort3(int[] arr) {
    if (arr[0] > arr[1]) {
        int temp = arr[0];
        arr[0] = arr[1];
        arr[1] = temp;
    }
    if (arr[1] > arr[2]) {
        int temp = arr[1];
        arr[1] = arr[2];
        arr[2] = temp;
    }
    if (arr[0] > arr[1]) {
        int temp = arr[0];
        arr[0] = arr[1];
        arr[1] = temp;
    }
}
```

```
    }
}
```

By following this process, we ensure that the array is sorted, regardless of the initial order of the elements.

8.6 Scaling Up: Sorting an Array of Size 4

We can extend the same approach to sort an array of size four. However, notice that once we have the largest number at the last position, the problem

reduces to sorting an array of size three. This realization will be crucial as we move forward.

```
public static void sort4(int[] arr) {
    for(int i = 0; i < 3; i++) {
        if (arr[i] > arr[i+1]) {
            int temp = arr[i];
            arr[i] = arr[i+1];
            arr[i+1] = temp;
        }
    }
    sort3(arr);
}
```

8.7 Recognizing the Pattern

As we incrementally solve larger versions of our sorting problem, we start to notice a pattern. For every array size, we first ensure the largest number ends up at the end of the array. Then, we apply the same sorting logic to the remaining elements in the array.

For instance, when sorting an array of size 5, we first find and move the largest number to the end of the array, then apply the sorting process for an array of size 4 to the remaining numbers:

```
public static void sort5(int[] arr) {
    for(int i = 0; i < 4; i++) {
        if (arr[i] > arr[i+1]) {
            int temp = arr[i];
            arr[i] = arr[i+1];
            arr[i+1] = temp;
```

```
    }
}
sort4(arr);
}
```

8.8 Generalizing the Algorithm

This pattern suggests a way to generalize our algorithm to sort an array of any size. For an array of size n , we first ensure the largest number ends up at the end, then recursively apply the same process to the remaining $n-1$ elements. By repeatedly applying this process, we ensure that the entire array ends up sorted.

```
public static void sort(int[] arr) {
    for(int i = arr.length - 1; i > 0; i--) {
        for(int j = 0; j < i; j++) {
            if (arr[j] > arr[j+1]) {
                int temp = arr[j];
                arr[j] = arr[j+1];
                arr[j+1] = temp;
            }
        }
    }
}
```

This algorithm, known as Bubble Sort, represents a simple yet effective solution to the sorting problem. It demonstrates how complex problems can often be broken down into simpler subproblems, providing a valuable lesson for problem-solving in computer science and beyond.

9 Comparable and Comparator

9.1 Sorting Complex Objects

In our previous lessons, we have primarily dealt with arrays of basic data types, such as integers, which are inherently easy to compare and sort. However, software development often involves sorting collections of complex objects, such as `Employee` or `Student` objects. These user-defined types lack a natural order for comparison, necessitating a custom solution.

9.2 Defining a Comparison Method

A rudimentary solution could involve appending a comparison method to our `Employee` or `Student` classes, like `compareSalary` or `compareGPA`. However, this design lacks versatility - each new attribute requires a fresh comparison method.

A superior solution employs the `Comparable` interface, comprising a single method, `compareTo`. This method, accepting an object of the same type, returns an integer indicating if the current object is lesser than, equal to, or greater than the input object.

The `Comparable` interface in Java is defined as follows:

```
public interface Comparable<T> {
    int compareTo(T o);
}
```

9.3 Implementing Comparable

Implementing the `Comparable` interface in any class endows it with sortable attributes. For example, sorting `Student` objects by GPA could be achieved by implementing `Comparable<Student>` in the `Student` class, defining `compareTo` to compare GPAs:

```
public class Student implements Comparable<Student> {
    private double gpa;
    // other fields and methods...
```

```

@Override
public int compareTo(Student other) {
    return Double.compare(this.gpa, other.gpa);
}
}

```

The array of `Student` objects can be sorted with our new comparison logic, `arr[j].compareTo(arr[j+1]) > 0`:

```

public static <T extends Comparable<T>> void sort(T[] arr) {
    for(int i = arr.length - 1; i > 0; i--) {
        for(int j = 0; j < i; j++) {
            if (arr[j].compareTo(arr[j+1]) > 0) {
                T temp = arr[j];
                arr[j] = arr[j+1];
                arr[j+1] = temp;
            }
        }
    }
}

```

9.4 Comparable in the Standard Library

The `Comparable` interface facilitates the creation of versatile sorting functions, capable of sorting arrays of any type implementing `Comparable`. This applies to user-defined types, as well as many built-in types in the Java standard library.

Implementing the `Comparable` interface communicates that a class is able to be compared with other instances of its type. This standard behavior empowers methods to interact with objects more abstractly and flexibly.

The Java standard library provides a `Comparable` interface similar to the one defined earlier. Many built-in classes such as `Integer`, `Double`, and `String` already implement this interface, enabling their comparison and sorting with no additional code. Our `Student` class can be modified to use `java.lang.Comparable` in place of our custom interface, maintaining the same `compareTo` method for use in our generic sort function.

9.5 Sorting Flexibility with Comparator

The `Comparable` interface is limited to a single `compareTo` method per class. When multiple sorting methods are required, Java's `Comparator` interface comes into play. `Comparator` represents different orderings for a specific type, allowing multiple comparators per class.

For instance, a `Comparator` for `Student` objects that orders by GPA could look like this:

```
import java
.util.Comparator;

public class StudentGPAComparator implements Comparator<Student> {
    @Override
    public int compare(Student s1, Student s2) {
        return Double.compare(s1.getGPA(), s2.getGPA());
    }
}
```

A modified version of our generic sort function, now accepting a `Comparator`, can compare elements of the array:

```
public static <T> void sort(T[] arr, Comparator<? super T> comparator) {
    for(int i = arr.length - 1; i > 0; i--) {
        for(int j = 0; j < i; j++) {
            if (comparator.compare(arr[j], arr[j+1]) > 0) {
                T temp = arr[j];
                arr[j] = arr[j+1];
                arr[j+1] = temp;
            }
        }
    }
}
```

9.6 Understanding Natural and Total Orderings

Implementing `Comparable` endows a class with a *natural ordering*, a default ordering used in operations like sorting. For example, `String` objects follow lexicographic order and `Integer` objects adhere to numerical order. When `Comparable<Student>` is implemented to compare GPAs, we establish GPA as the natural ordering for `Student` objects.

Java's `Comparator` interface allows the definition of additional orderings, known as *total orderings*. While the natural ordering is default, total orderings, which also follow rules of completeness, transitivity, and antisymmetry, provide flexibility to define any suitable ordering.

The `Comparable` and `Comparator` interfaces epitomize the power of abstraction in computer science. By concentrating on the core concept of order and essential operations involving order, we can create general, flexible, and reusable code to work with a broad spectrum of data types and orderings.

In the next chapter, we will delve into how these concepts are employed in data structures like trees and heaps to manage data for efficient searching and sorting. For now, contemplate the elegance and versatility of the `Comparable` and `Comparator` interfaces, and the compelling concept of order they embody.

10 Bubble, Selection, and Insertion Sort

10.1 Bubble Sort

Bubble sort is a simple sorting algorithm that works by repeatedly comparing and swapping adjacent elements in an array until the array is sorted. Bubble sort is easy to implement and understand, but it is not very efficient for large arrays. The name bubble sort comes from the idea that the larger elements “bubble up” to the end of the array after each iteration.

10.1.1 Pseudocode

The pseudocode for bubble sort is as follows:

```
procedure bubbleSort(A : list of sortable items)
    n = length(A)
    repeat
        swapped = false
        for i = 1 to n-1 inclusive do
            if A[i-1] > A[i] then
                swap(A[i-1], A[i])
                swapped = true
            end if
        end for
        n = n - 1
    until not swapped
end procedure
```

The procedure takes an array A of sortable items and sorts it in ascending order. The procedure uses a variable n to keep track of the unsorted part of the array, and a variable swapped to indicate whether any swaps occurred in the current iteration. The procedure repeats until no swaps are made, which means the array is sorted.

10.1.2 An Example

1. Let us see an example of bubble sort on the following array:

```
[5, 1, 4, 2, 8]
```

2. We start with `n = 5` and `swapped = false`. We compare the first and second elements, 5 and 1, and swap them since $5 > 1$. We set `swapped` to true to indicate that a swap occurred. We compare the second and third elements, 5 and 4, and swap them as well. We continue to compare and swap the adjacent elements until we reach the end of the array. The array after the first iteration is:

```
[1, 4, 2, 5, 8]
```

3. We decrement `n` by 1, since the last element 8 is now in its correct position. We repeat the process with `n = 4` and `swapped = false`. We compare and swap the first and second elements, 1 and 4, since $1 < 4$. We compare and swap the second and third elements, 4 and 2, since $4 > 2$. We compare the third and fourth elements, 5 and 8, but do not swap them since $5 < 8$. The array after the second iteration is:

```
[1, 2, 4, 5, 8]
```

4. We decrement `n` by 1, since the last element 5 is now in its correct position. We repeat the process with `n = 3` and `swapped = false`. We compare and swap the first and second elements, 1 and 2, since $1 < 2$. We compare and swap the second and third elements, 2 and 4, since $2 < 4$. The array after the third iteration is:

```
[1, 2, 4, 5, 8]
```

5. We decrement `n` by 1, since the last element 4 is now in its correct position. We repeat the process with `n = 2` and `swapped = false`. We compare and swap the first and second elements, 1 and 2, since $1 < 2$. The array after the fourth iteration is:

```
[1, 2, 4, 5, 8]
```

6. We decrement `n` by 1, since the last element 2 is now in its correct position. We repeat the process with `n = 1` and `swapped = false`. We do not compare any elements, since there is only one element left. The array after the fifth iteration is:

```
[1, 2, 4, 5, 8]
```

7. We exit the loop, since `swapped` is false, which means the array is sorted.

10.1.3 Lab Instructions

1. The template uses Java generics to create a generic class `Main` that can sort arrays of any type that implements the `Comparable` interface. Generics are a way of implementing generic programming in Java, which allows you to write code that can work with different types of objects without casting or risking `ClassCastException`.
2. The constructor of the `Main` class takes an array of type `T` as a parameter and assigns it to the `data` field. The `data` field is also of type `T`, which means it can store any type of object that implements `Comparable`.
3. The `BubbleSort` method returns an array of type `T` that is sorted in ascending order using the bubble sort algorithm. The method uses a local variable `sorted` to store a copy of the `data` array and then modifies it using the pseudocode provided in the assignment instructions.
4. To compare and swap the elements of the `sorted` array, you need to use the `compareTo` method of the `Comparable` interface. The `compareTo` method returns a negative integer, zero, or a positive integer if the current object is less than, equal to, or greater than the specified object. For example, if you want to compare the elements at index `i-1` and `i`, you can write:

```
if (sorted[i-1].compareTo(sorted[i]) > 0) {  
    // swap the elements  
}
```

5. To swap the elements of the `sorted` array, you can use a temporary variable of type `T` to store one of the elements, and then assign the other element to its place. For example, if you want to swap the elements at index `i-1` and `i`, you can write:

```
T temp = sorted[i-1];  
sorted[i-1] = sorted[i];  
sorted[i] = temp;
```

6. After performing the bubble sort algorithm, the method returns the sorted array.
7. To test your code, you can create an object of the `Main` class with different types of arrays, such as `Integer`, `String`, or `Double`, and call the `BubbleSort` method on them. You can print the original and sorted arrays to check the output. For example, you can write:

```
Integer[] intArray = {5, 1, 4, 2, 8};  
Main<Integer> intMain = new Main<>(intArray);  
Integer[] intSorted = intMain.BubbleSort();
```

```
System.out.println("Original array: " + Arrays.toString(intArray));
System.out.println("Sorted array: " + Arrays.toString(intSorted));
```

The output should be:

```
Original array: [5, 1, 4, 2, 8]
Sorted array: [1, 2, 4, 5, 8]
```

10.2 Selection Sort

10.2.1 Introduction

Selection sort is a sorting algorithm that repeatedly finds the smallest element in the unsorted part of the array and swaps it with the first element of the unsorted part. This way, the array is divided into two subarrays: one that contains the sorted elements and one that contains the unsorted elements. The process is repeated until the entire array is sorted. Selection sort is a simple and efficient algorithm that works well for small arrays, but it has a high time complexity of $O(n^2)$ for large arrays.

10.2.2 Pseudocode

The pseudocode for selection sort is as follows:

```
Procedure selection_sort(array, N)
    array - array of items to be sorted
    N - size of array
Begin
    For i = 1 to N-1
        Begin
            Set min = i
            For j = i+1 to N
                Begin
                    If array[j] < array[min] Then
                        min = j
                    End If
                End For
                // Swap the minimum element with the current element
                If min != i Then
                    Swap array[min] and array[i]
```

```
    End If
End For
End Procedure
```

10.2.3 An Example

1. Let us see an example of how selection sort works on the following array:

```
[64, 25, 12, 22, 11]
```

2. In the first iteration, the algorithm finds the smallest element (11) in the unsorted part of the array and swaps it with the first element of the unsorted part (64). The array becomes:

```
[11, 25, 12, 22, 64]
```

3. In the second iteration, the algorithm finds the smallest element (12) in the remaining unsorted part of the array and swaps it with the second element of the unsorted part (25). The array becomes:

```
[11, 12, 25, 22, 64]
```

4. In the third iteration, the algorithm finds the smallest element (22) in the remaining unsorted part of the array and swaps it with the third element of the unsorted part (25). The array becomes:

```
[11, 12, 22, 25, 64]
```

5. In the fourth iteration, the algorithm finds the smallest element (25) in the remaining unsorted part of the array and swaps it with the fourth element of the unsorted part (25). The array becomes:

```
[11, 12, 22, 25, 64]
```

6. In the fifth iteration, the algorithm finds the smallest element (64) in the remaining unsorted part of the array and swaps it with the fifth element of the unsorted part (64). The array becomes:

```
[11, 12, 22, 25, 64]
```

7. The array is now sorted and the algorithm terminates.

10.2.4 Instructions

1. The template is a generic class that can sort any type of data that implements the `Comparable` interface. This means that the data type `T` must have a method `compareTo(T other)` that returns a negative, zero, or positive integer depending on whether the current object is less than, equal to, or greater than the other object.
2. The constructor of the class takes an array of type `T` as a parameter and assigns it to the data field. The data field is the array that needs to be sorted.
3. The `SelectionSort()` method is where the students need

to write the code for the selection sort algorithm. The method should return a sorted array of type `T`.

4. The students can use the pseudocode provided in the previous section as a guide for writing the code. They need to use a for loop to iterate over the unsorted part of the array, find the minimum element, and swap it with the first element of the unsorted part. They can use the `compareTo()` method to compare the elements of the array. They can use a temporary variable to store the value of the element to be swapped.
5. The students can test their code by creating an object of the `Main` class with different types of data, such as integers, strings, or custom objects, and calling the `SelectionSort()` method on it. They can print the original and sorted arrays to check the output. They can also use different sizes of arrays to see how the algorithm performs.

10.3 Insertion Sort

10.3.1 Introduction

Insertion sort is a simple and adaptive sorting technique that works by inserting each element into its correct position in a sorted subarray. The sorted subarray starts with the first element and grows by one element in each iteration. The element to be inserted is compared with the previous elements in the sorted subarray and swapped if they are larger. This process continues until the element finds its correct position or reaches the beginning of the array.

10.3.2 Pseudocode

The pseudocode for insertion sort is as follows:

```

procedure insertionSort (A: list of sortable items)
    n = length (A)
    for i = 1 to n - 1 do
        j = i
        while j > 0 and A [j-1] > A [j] do
            swap (A [j], A [j-1])
            j = j - 1
        end while
    end for
end procedure

```

10.3.3 An Example

Consider the array: [3, 5, 7, 11, 13, 2, 9, 6]

Below is a detailed walkthrough of each iteration:

1. Iteration 1: [3, 5, 7, 11, 13, 2, 9, 6]

The sorted subarray is [3] and the element to be inserted is 5. Since 5 is larger than 3, no swap is needed.

2. Iteration 2: [3, 5, 7, 11, 13, 2, 9, 6]

The sorted subarray is [3, 5] and the element to be inserted is 7. Since 7 is larger than 5, no swap is needed.

3. Iteration 3: [3, 5, 7, 11, 13, 2, 9, 6]

The sorted subarray is [3, 5, 7] and the element to be inserted is 11. Since 11 is larger than 7, no swap is needed.

4. Iteration 4: [3, 5, 7, 11, 13, 2, 9, 6]

The sorted subarray is [3, 5, 7, 11] and the element to be inserted is 13. Since 13 is larger than 11, no swap is needed.

5. Iteration 5: [3, 5, 7, 11, 13, 2, 9, 6]

The sorted subarray is [3, 5, 7, 11, 13] and the element to be inserted is 2. Since 2 is smaller than 13, 13 and 2 are swapped.

After swapping: [3, 5, 7, 11, 2, 13, 9, 6]

6. Iteration 6: [3, 5, 7, 11, 2, 13, 9, 6]

The sorted subarray is [3, 5, 7, 11, 2, 13] and the element to be inserted is 2. Since 2 is smaller than 11, 11 and 2 are swapped.

After swapping: [3, 5, 7, 2, 11, 13, 9, 6]

7. Iteration 7: [3, 5, 7,

2, 11, 13, 9, 6]`

The sorted subarray is [3, 5, 7, 2, 11, 13] and the element to be inserted is 2. Since 2 is smaller than 7, 7 and 2 are swapped.

After swapping: [3, 5, 2, 7, 11, 13, 9, 6]

8. Iteration 8: [3, 5, 2, 7, 11, 13, 9, 6]

The sorted subarray is [3, 5, 2, 7, 11, 13] and the element to be inserted is 2. Since 2 is smaller than 5, 5 and 2 are swapped.

After swapping: [3, 2, 5, 7, 11, 13, 9, 6]

9. Iteration 9: [3, 2, 5, 7, 11, 13, 9, 6]

The sorted subarray is [3, 2, 5, 7, 11, 13] and the element to be inserted is 2. Since 2 is smaller than 3, 3 and 2 are swapped.

After swapping: [2, 3, 5, 7, 11, 13, 9, 6]

10. Iteration 10: [2, 3, 5, 7, 11, 13, 9, 6]

The sorted subarray is [2, 3, 5, 7, 11, 13] and the element to be inserted is 9. Since 9 is smaller than 13, 13 and 9 are swapped.

After swapping: [2, 3, 5, 7, 11, 9, 13, 6]

11. Iteration 11: [2, 3, 5, 7, 11, 9, 13, 6]

The sorted subarray is [2, 3, 5, 7, 11, 9, 13] and the element to be inserted is 9. Since 9 is smaller than 11, 11 and 9 are swapped.

After swapping: [2, 3, 5, 7, 9, 11, 13, 6]

12. Iteration 12: [2, 3, 5, 7, 9, 11, 13, 6]

The sorted subarray is [2, 3, 5, 7, 9, 11, 13] and the element to be inserted is 9. Since 9 is larger than 7, no swap is needed.

13. Iteration 13: [2, 3, 5, 7, 9, 11, 13, 6]

The sorted subarray is [2, 3, 5, 7, 9, 11, 13] and the element to be inserted is 6. Since 6 is smaller than 13, 13 and 6 are swapped.

After swapping: [2, 3, 5, 7, 9, 11, 6, 13]

14. **Iteration

14**: [2, 3, 5, 7, 9, 11, 6, 13]

The sorted subarray is [2, 3, 5, 7, 9, 11, 6, 13] and the element to be inserted is 6. Since 6 is smaller than 11, 11 and 6 are swapped.

After swapping: `[2, 3, 5, 7, 9, 6, 11, 13]`

15. Iteration 15: [2, 3, 5, 7, 9, 6, 11, 13]

The sorted subarray is [2, 3, 5, 7, 9, 6, 11, 13] and the element to be inserted is 6. Since 6 is smaller than 9, 9 and 6 are swapped.

After swapping: [2, 3, 5, 7, 6, 9, 11, 13]

16. Iteration 16: [2, 3, 5, 7, 6, 9, 11, 13]

The sorted subarray is [2, 3, 5, 7, 6, 9, 11, 13] and the element to be inserted is 6. Since 6 is smaller than 7, 7 and 6 are swapped.

After swapping: [2, 3, 5, 6, 7, 9, 11, 13]

17. Iteration 17: [2, 3, 5, 6, 7, 9, 11, 13]

The sorted subarray is [2, 3, 5, 6, 7, 9, 11, 13] and the element to be inserted is 6. Since 6 is larger than 5, no swap is needed.

Final array: [2, 3, 5, 6, 7, 9, 11, 13]

The sorted subarray is [2, 3, 5, 6, 7, 9, 11, 13] and the array is fully sorted.

10.3.4 Instructions

1. Start by understanding how insertion sort works and what it does to sort an array.
2. Study the given template for the Main class and understand how it uses Java generics to sort an array of any type that implements the Comparable interface.
3. Implement the `InsertionSort` method by following the pseudocode provided above. Students should create a copy of the input array and modify the copy in-place to produce the sorted array.
4. Test their implementation using a variety of input arrays and verify that the output is correct.
5. Encourage students to optimize their implementation by considering edge cases and identifying areas for improvement.
6. Remind students to comment their code and explain their thought process as they work through the assignment. This will help them develop their coding skills and become better problem solvers.

11 Linear Search

Linear search is one of the simplest searching algorithms. The linear search algorithm scans one item at a time, without jumping to any item. It starts at the beginning of a list (or array) and sequentially checks each element until it finds a match.

Here's a general breakdown of how the linear search algorithm works:

1. It starts at the first item in the list.
2. It moves from item to item, checking if the current item is equal to the item being searched.
3. If a match is found, the algorithm stops and returns the index of the item.
4. If the algorithm reaches the end of the list without finding a match, it returns a signal that no match was found (often this is -1 or null).

Linear search can be applied to any type of data. However, since it searches element-by-element, it can be inefficient for large lists or arrays.

11.1 Pseudocode

The pseudocode for linear search is straightforward:

```
procedure LinearSearch(A, key)
    for i from 1 to length(A)
        if A[i] equals key
            return i
    return -1
```

This pseudocode represents a function called `LinearSearch` that takes two parameters: a list (or array) `A` and a `key` to search for in `A`. The function returns the index of `key` if found, otherwise it returns -1.

11.2 An Example

Let's consider the following example: We have an array [10, 15, 20, 25, 30, 35] and we want to find the number 20.

1. **Iteration 1:** The first element is 10. 10 is not equal to 20, so we move on to the next element.
2. **Iteration 2:** The second element is 15. 15 is not equal to 20, so we move on to the next element.
3. **Iteration 3:** The third element is 20. 20 is equal to 20, so we stop searching. The index of the number 20 is 2 (considering the first index as 0).

Therefore, the linear search algorithm would return 2 as the index of the number 20 in the given array.

11.3 Instructions

Here are some instructions to help you start working on the assignment:

1. The template uses Java generics to create a generic class Main that can search for a key element in an array of any type that implements the Comparable interface. Generics are a way of implementing generic programming in Java, which allows you to write code that can work with different types of objects without casting or risking ClassCastException.
2. The constructor of the Main class takes an array of type T as a parameter and assigns it to the array field. The array field is also of type T, which means it can store any type of object that implements Comparable.
3. The `LinearSearch` method returns an int that is the index of the key element in the array, or -1 if the key element is not found. The method uses a for loop to iterate over the array and compare each element with the key using the `compareTo` method of the Comparable interface. The `compareTo` method returns a negative integer, zero, or a positive integer if the current object is less than, equal to, or greater than the specified object. For example, if you want to compare the element at index i with the key, you can write:

```
if (array[i].compareTo(key) == 0) {  
    // return the index  
}
```

The `LinearSearch` method has a time complexity of $O(n)$, where

n is the size of the array, because it may have to scan the entire array to find the key element.

4. To test your code, you can create an object of the Main class with different types of arrays, such as Integer, String, or Double, and call the `LinearSearch` method on them. You can print the key element and the index returned by the methods to check the output. For example, you can write:

```
```java
Integer[] intArray = {1, 2, 4, 5, 8};
Main<Integer> intMain = new Main<>(intArray);
Integer key = 4;
int linearIndex = intMain.LinearSearch(key);
System.out.println("Key element: " + key);
System.out.println("Linear search index: " + linearIndex);
```

```

The output of the above code should be:

```
```
Key element: 4
Linear search index: 2
```

```

Part V

Sequential Collections of Data

12 Collections of Data

12.1 (Active) Memory in Computers

To fully grasp the concept of data collections in Java, we must first lay a foundational understanding of memory in computing systems. It is helpful to visualize the computer's memory as a vast grid, with each cell in the grid being capable of storing a certain number of bits. This simple grid becomes the underpinning infrastructure for data storage, holding variables that serve a multitude of functions in our programming ventures.

| | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 1 | 1 | 0 | 1 | 0 | 1 | 0 | 1 | 0 | 1 |
| 0 | 1 | 0 | 1 | 1 | 0 | 1 | 0 | 1 | 0 | 1 | 0 |
| 1 | 0 | 1 | 1 | 0 | 1 | 0 | 1 | 0 | 1 | 0 | 1 |
| 0 | 1 | 0 | 1 | 1 | 0 | 1 | 0 | 1 | 0 | 1 | 0 |

Figure 12.1: A diagram showing a grid of memory cells, each capable of storing a single bit.

12.2 The Concept of Data Collections

In programming, the need often arises to handle groups of similar variables. These groups are what we refer to as data collections. A data collection is essentially an assembly of bits in the

memory, representing a group of variables. For instance, you could have a collection of four 2-bit variables, which together would take up eight bits of memory. But the behavior of these collections can vary, particularly when considering size constraints.

12.3 Types of Collections: Fixed and Dynamic

When analyzing the size of data collections, we encounter two main types: fixed (or static) and dynamic collections.

Fixed or Static Collections: As the name suggests, in a static collection, the number of variables is constant. This implies that the collection's size is immutable, and the total memory occupancy remains consistent. An array of 10 integers is a perfect example of a static collection - regardless of the values stored in it, the array always contains ten integers.

| Data | Data | Data | | | | | | | | | |
|------|------|------|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 1 | 1 | 1 | 0 | 0 | 1 | 1 | 0 | 1 |
| 1 | 1 | 1 | 0 | 1 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |

Figure 12.2: A diagram showing a static collection of data where the five 4-bit variables are placed sequentially in memory.

Dynamic Collections: Dynamic collections differ from their static counterparts in that their sizes are subject to change over time. Consider, for example, a collection storing the names of your favorite musicians. As your taste in music evolves, so does this list, reflecting an expansion or contraction based on your preferences. Thus, it is a dynamic collection.

12.4 Memory Layout of Collections

Let's take a deeper look into how these collections are stored in memory. Note that the memory allocation of these collections can be influenced by factors such as the variables' data type and size.

A common approach to store a collection of two 4-bit numbers is sequentially, positioning them back-to-back. However, this is not the only strategy. Depending on the system's memory management and data alignment methodologies, 'padding' could be introduced, which involves inserting extra bits between variables to align them properly in memory.

Regardless of these variations, a crucial consideration when working with collections is devising a mechanism to locate the next variable in the sequence. This task is straightforward for static collections due to the known size of each variable. But how about dynamic collections?

12.5 Tackling Dynamic Collections

Dynamic collections present an interesting challenge due to their mutable size. One way to navigate this issue is by earmarking a portion of the memory to store the length of the collection. For instance, the first four bits could be utilized to denote the size of the collection.

By adopting this method, we can efficiently locate subsequent data elements in our collection and discern the collection's end point. Bear in mind that the data in collections doesn't have to be stored adjacently. We can opt to pad each element with a fixed number of bits. As long as the padding size remains consistent, we can still locate the next piece of data.

12.6 Memory Layout and Program Performance

The manner in which data collections are stored in memory can have a significant impact on the performance of a program. For instance, accessing memory sequentially (in a pattern that matches the memory layout of the collection) allows the system to leverage cache lines, which are blocks of memory that are read into the CPU's cache. Due to the way modern CPUs are designed, once a part of the memory is read, an entire cache line (usually 64 or 128 bytes) is loaded into the cache. Thus, sequential access often results in fewer cache misses, which improves the performance of your program.

In contrast, random access to memory can lead to frequent cache misses, as each access may require loading a different cache line into the CPU cache, slowing down your program. Therefore, understanding the memory layout of your data collections and designing your program to access data in a manner that complements this layout can lead to significant performance improvements.

| Length | Data | Pad | Data | Pad |
|--------|------|-----|------|-----|
| 0 | 0 | 1 | 0 | 0 |
| 1 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 |

Figure 12.3: A diagram showing a dynamic collection of data where the first 4 bits represent the length of the collection, and the next few bits are the collection of variables. Each variable will be separated by 2-bit padding.

So, as you delve deeper into the world of data collections, remember that your approach to handling memory can either elevate or hinder your program's efficiency. The choice, as always, is in your hands!

12.7 Summary and Conclusion

This chapter provided a comprehensive overview of data collections in programming, focusing mainly on their storage in computer memory. It began by explaining the fundamental understanding of memory in computing systems, portraying it as a vast grid where each cell stores a certain number of bits. This concept helped set the groundwork for understanding how data collections, which are groups of similar variables, are stored in memory.

The chapter then discussed two main types of data collections: fixed (static) and dynamic collections. While fixed collections have a constant number of variables and thus a constant memory size, dynamic collections have a variable size that can change over time. An illustration was given of both types to aid understanding.

The memory layout of these collections was then delved into, noting that how these collections are stored in memory can be influenced by the variables' data type and size. The chapter also explained the strategies for accessing variables in both static and dynamic collections.

Moreover, the chapter emphasized the significant impact of the memory layout of data collections on program performance. Sequential memory access in line with the memory layout could leverage cache lines and enhance the program's performance, while random access could lead to frequent cache misses and slow down the program.

In conclusion, understanding data collections, how they are stored in memory, and how to access this data efficiently are critical aspects of programming. By considering these factors, programmers can significantly influence the performance of their programs. Thus, the mastery of data collections is not only essential for writing efficient code but is also a vital skill for optimizing the overall performance of a program.

13 Operations on Lists

In the prior lecture, we journeyed through the concept of collections of data and their manifestations in computer memory. Such collections were referred to as lists, a data structure central to understanding the organization and manipulation of data in programming. Our objective is to delve into the practical implementation of these lists in Java. However, before we embark on that, we must consider the types of operations we desire our lists to perform. These operations guide our implementation strategy and ensure that our list behaves in the ways we want it to. In this chapter, we'll be investigating operations performed on a List Abstract Data Type (ADT).

13.1 Adding Data to a List

As you may recall, a list is a collection of data elements. A key operation, therefore, is the ability to add data to this collection. Let's visualize a simple list:

| | | | | | |
|-------|---|---|---|---|---|
| 1 | 2 | 3 | 4 | 5 | 6 |
| <hr/> | | | | | |

There are multiple ways to insert new data into this list. For instance, you might want to add a new element, say “X”, at the end:

| | | | | | | |
|-------|---|---|---|---|---|---|
| 1 | 2 | 3 | 4 | 5 | 6 | X |
| <hr/> | | | | | | |

Alternatively, you could place “X” at the beginning:

| | | | | | | |
|-------|---|---|---|---|---|---|
| X | 1 | 2 | 3 | 4 | 5 | 6 |
| <hr/> | | | | | | |

Or perhaps insert “X” somewhere in the middle:

| | | | | | | |
|-------|---|---|---|---|---|---|
| 1 | 2 | 3 | X | 4 | 5 | 6 |
| <hr/> | | | | | | |

There are three distinct methods to add data to our list, and each method may serve different needs in your program. Let's discuss how these operations are implemented in the List ADT provided by the Java standard library, `java.util.List`.

13.1.1 Appending an Element

Appending an element to the end of a list is perhaps the most straightforward method of adding data. The `add` method provided by `java.util.List` is used for this operation. Here is how you might use it:

```
listObject.add(itemToAdd);
```

13.1.2 Prepending an Element

Inserting an element at the beginning of a list is also a common operation, often called “prepend”. Though the `java.util.List` interface doesn't provide a `prepend` method, we can use the overloaded `add` method, which accepts an index and an element. To prepend, we simply pass `0` as the index:

```
listObject.add(0, itemToAdd);
```

13.1.3 Inserting an Element at a Specific Index

As hinted in the prepend operation, the `add` method in `java.util.List` allows for adding an element at any index in the list:

```
listObject.add(index, itemToAdd);
```

13.1.4 Adding All Elements from Another List

Sometimes, you may want to merge one list into another. The `addAll` method allows you to add all the elements from another list to the end of the current list:

```
listObject.addAll(anotherList);
```

13.2 The Role of Abstract Data Types (ADTs)

We've mentioned the term "Abstract Data Type" (ADT) several times so far, but what does it really mean? An ADT is a high-level description of a collection of data and the operations that can be performed on that data. It is "abstract" in that it describes what operations are to be done but not how these operations will be implemented.

In the context of lists,

the List ADT defines a list's behavior. For instance, we can append, prepend, or insert an element at a specific index, but we don't care about how these operations are performed.

Why do we need ADTs? ADTs allow us to abstract away the details of the data structure's implementation. They encapsulate the data and provide a well-defined interface to interact with it, ensuring that data remains consistent and operations on the data are predictable. Moreover, by using ADTs, different implementations of the same type of data structure can be swapped seamlessly, without changing the code that uses the data structure.

In your journey with Java so far, you've only used one kind of list – the ArrayList. As you progress, you'll encounter other types of lists, like LinkedLists and Stacks, each with its own strengths, weaknesses, and suitable use cases. All these list types adhere to the same List ADT, allowing us to switch from one type to another depending on our requirements, while our code remains largely the same.

In the next sections of this course, we will dive deeper into different types of lists and their unique characteristics. For now, the important takeaway is that ADTs allow us to focus on *what* operations we want to perform without worrying about *how* they are implemented.

13.2.1 Summary of Addition Operations

Here is a summary of the adding operations available in the `java.util.List` ADT:

| Modifier and Type | Method and Description |
|-------------------|--|
| boolean | <code>add(E e)</code> Appends the specified element to the end of this list (optional operation). |
| void | <code>add(int index, E element)</code> Inserts the specified element at the specified position in this list (optional operation). |
| boolean | <code>addAll(Collection<? extends E> c)</code> Appends all of the elements in the specified collection to the end of this list, in the order that they are returned by the specified collection's iterator (optional operation). |

| Modifier and Type | Method and Description |
|-------------------|--|
| boolean | <code>addAll(int index, Collection<? extends E> c)</code> Inserts all of the elements in the specified collection into this list at the specified position (optional operation). |

13.3 Removing data from our List

13.4 Removing Data from a List

Having explored the addition of data to a list, let's now shift our attention to its counterpart: removing data from the list. Remember our list from the previous section?

| | | | | | |
|---|---|---|---|---|---|
| 1 | 2 | 3 | 4 | 5 | 6 |
| | | | | | |

Let's consider different ways to remove data from this list.

13.4.1 Removing an Element at a Specific Index

Sometimes, you might need to remove an item from a specific position within the list. Java's List ADT provides the `remove(int index)` method for this purpose. This method removes the element at the specified position:

```
listObject.remove(index);
```

13.4.2 Removing the First Occurrence of an Element

At times, you might not know (or care about) the index of the element you want to remove, but you know the value of the element. In such cases, you can use the `remove(Object o)` method, which removes the first occurrence of the specified element from the list:

```
listObject.remove(objectToRemove);
```

13.4.3 Removing All Elements from Another List

Consider that you have two lists, and you want to remove all elements in the second list from the first one. You can use the `removeAll(Collection<?> c)` method, which removes from the current list all of its elements that are contained in the specified collection:

```
listObject.removeAll(anotherList);
```

13.4.4 Removing All Elements from the List

In some scenarios, you might want to clear your list entirely. The `clear()` method comes in handy for this, as it removes all elements from the list:

```
listObject.clear();
```

Like the addition operations, these removal operations offer different ways to manage the data in our list, providing flexibility based on the specific needs of our program.

13.4.5 Summary of Removal Operations

Here is a summary of the removal operations available in the `java.util.List` ADT:

| Modifier and Type | Method and Description |
|-------------------|---|
| E | <code>remove(int index)</code> Removes the element at the specified position in this list (optional operation). |
| boolean | <code>remove(Object o)</code> Removes the first occurrence of the specified element from this list, if it is present (optional operation). |
| boolean | <code>removeAll(Collection<?> c)</code> Removes from this list all of its elements that are contained in the specified collection (optional operation). |
| void | <code>clear()</code> Removes all of the elements from this list (optional operation). |

13.5 Searching for data in our List

13.6 Searching in a List

Beyond just adding and removing data from our list, we often need to find or check for the existence of certain elements in our list. Java's List ADT provides several methods for such search operations. Let's dive into these operations using our familiar list:

| | | | | | |
|---|---|---|---|---|---|
| 1 | 2 | 3 | 4 | 5 | 6 |
|---|---|---|---|---|---|

13.6.1 Checking If an Element Exists in the List

Sometimes, all we need to know is whether a certain element exists in our list. The `contains(Object o)` method serves this purpose, returning `true` if this list contains the specified element, and `false` otherwise:

```
boolean contains = listObject.contains(objectToCheck);
```

13.6.2 Checking If All Elements of Another Collection Exist in the List

If you have a collection of elements and you want to check whether all these elements exist in your list, you can use the `containsAll(Collection<?> c)` method. It returns `true` if the list contains all of the elements of the specified collection:

```
boolean containsAll = listObject.containsAll(anotherCollection);
```

13.6.3 Finding the Index of an Element

If you want to find the position of a certain element in your list, Java provides the `indexOf(Object o)` method. This method returns the index of the first occurrence of the specified element in this list, or -1 if this list does not contain the element:

```
int index = listObject.indexOf(objectToFind);
```

13.6.4 Finding the Last Index of an Element

In a list with duplicate elements, you might be interested in finding the last occurrence of an element. In this case, you can use the `lastIndexOf(Object o)` method, which returns the index of the last occurrence of the specified element in this list, or -1 if this list does not contain the element:

```
int lastIndex = listObject.lastIndexOf(objectToFind);
```

13.6.5 Summary of Search Operations

Here is a summary of the search operations available in the `java.util.List` ADT:

| Modifier and Type | Method and Description |
|-------------------|---|
| boolean | <code>contains(Object o)</code> Returns true if this list contains the specified element. |
| boolean | <code>containsAll(Collection<?> c)</code> Returns true if this list contains all of the elements of the specified collection. |
| int | <code>indexOf(Object o)</code> Returns the index of the first occurrence of the specified element in this list, or -1 if this list does not contain the element. |
| int | <code>lastIndexOf(Object o)</code> Returns the index of the last occurrence of the specified element in this list, or -1 if this list does not contain the element. |

Search operations are a fundamental part of manipulating lists. They allow us to locate and verify the presence of data, enhancing our ability to interact with our list and derive meaningful results.

13.7 Miscellaneous Operations on a List

In addition to adding, removing, and searching elements in our list, Java's List ADT provides a variety of other useful operations. These allow us to manipulate and inquire our list in a more sophisticated manner.

13.7.1 Accessing an Element

Sometimes, we need to retrieve the element at a specific position in our list without removing it. The `get(int index)` method provides this functionality, returning the element at the specified position:

```
E element = listObject.get(index);
```

13.7.2 Modifying an Element

What if we need to change an element at a specific position? Java provides the `set(int index, E element)` method. This method replaces the element at the specified position in this list with the specified element:

```
listObject.set(index, newElement);
```

13.7.3 Determining the Size of the List

When working with lists, it's often necessary to know the number of elements present. The `size()` method returns the number of elements in this list:

```
int size = listObject.size();
```

13.7.4 Converting the List to an Array

On occasion, we may need to convert our list into an array. The `toArray()` method fulfills this purpose, returning an array containing all the elements in this list in proper sequence:

```
Object[] array = listObject.toArray();
```

13.7.5 Summary of Miscellaneous Operations

Here is a summary of the miscellaneous operations available in the `java.util.List` ADT:

| Modifier and Type | Method and Description |
|-------------------|---|
| E | <code>get(int index)</code> Returns the element at the specified position in this list. |

| Modifier and Type | Method and Description |
|-----------------------|---|
| <code>E</code> | <code>set(int index, E element)</code> Replaces the element at the specified position in this list with the specified element (optional operation). |
| <code>int</code> | <code>size()</code> Returns the number of elements in this list. |
| <code>Object[]</code> | <code>toArray()</code> Returns an array containing all of the elements in this list in proper sequence (from first to last element). |

These miscellaneous operations on a list help us to manipulate and utilize our lists effectively. In the subsequent sections, we will delve deeper into different types of lists and how these operations can be applied differently. Stay tuned!

13.8 Summary and Conclusion

This chapter aimed to offer a comprehensive understanding of operations that can be performed on lists using the Java programming language. It introduced key manipulations including adding, removing, and searching for elements in a list, as well as several other useful list operations.

We began by delving into the process of adding elements to a list, where we considered various insertion methods such as appending, prepending, and inserting at a specific index. We also explored the ability to add all elements from another list into the current list. These addition operations provide flexibility and adaptability in managing data in our list as required by the program.

Next, we moved on to the concept of removing elements from a list. We investigated removing elements at specific indices, removing the first occurrence of an element, and removing all elements from a list or all elements present in another list. These removal operations are just as vital in data management, providing diverse ways to regulate the contents of our list.

Subsequently, we analyzed search operations within a list. These operations are crucial in data retrieval, determining the existence and position of data elements in our list. Methods like `contains(Object o)`, `containsAll(Collection<?> c)`, `indexOf(Object o)`, and `lastIndexOf(Object o)` were discussed in depth.

Furthermore, we examined miscellaneous operations that provide additional functionality in our list. These include accessing and modifying elements at specific indices, determining the size of the list, and converting the list into an array.

Throughout our exploration, we highlighted the significant role of Abstract Data Types (ADTs). By standardizing the operations that can be performed on a list, the List ADT allows us to focus on what we want to accomplish without concerning ourselves with the underlying implementation details. This abstraction brings forth code consistency, reliability, and the potential for seamless integration of different list implementations.

In conclusion, understanding and utilizing list operations in Java is fundamental for effective data management in programming. The breadth of operations available affords us immense flexibility in manipulating lists to align with our program's requirements. As we proceed in this course, we will delve deeper into the different types of lists and their unique characteristics, all the while, leveraging the power and versatility of the List ADT.

14 Implementing Array-backed Lists

15 Implementing Linked Lists

In this chapter, we'll take a closer look at how objects interact in Java. To fully understand this interaction, we first need to grasp the concept of references in Java, memory allocation for Java objects, and how different objects can interact through references. Later, we'll also examine how these principles apply to implementing linked lists in Java.

15.1 Understanding References in Java

In Java programming, the statement that “all objects are references” is a cornerstone. Understanding what this statement means is a fundamental step in mastering Java.

In Java, when we say an object is a *reference*, it means that the object acts as an address pointer, referencing a specific location in memory. This location is where the actual data associated with that object is stored.

Consider your computer’s memory as a large grid, as illustrated in Table 6.1. Each cell can store some data.

Table 15.1: Representation of computer’s memory

| | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

For any object in Java, it’s not the data of the object that is directly stored in the variable, but rather the *reference* to the data.

15.2 Memory Allocation for Java Objects

Consider the following example, in which we have a `Dog` class.

```
public class Dog {  
    String name;  
    int age;  
}
```

Initially, the `Dog` class doesn't occupy any memory space. It simply serves as a blueprint for creating `Dog` objects.

Now, let's create a new instance of `Dog` and assign it to the variable `myDog`:

```
Dog myDog = new Dog();
```

This operation performs two main actions:

1. Memory is allocated for a new `Dog` object.
2. The address of that memory is stored in the `myDog` variable.




Figure 15.1: Dog Object Memory Allocation

So, `myDog` is essentially a pointer that points to the location in memory where the `Dog` object's data is stored.

15.3 Objects Interacting Through References

Now, let's consider a scenario where we have two classes, `Dog` and `HomeOwner`. The `Dog` class is the same as before, but the `HomeOwner` class is a bit more complex.

```
public class HomeOwner {  
    String name;  
    Dog pet;  
}
```

Here, the `HomeOwner` class has a `pet` field of type `Dog`. This means a `HomeOwner` can own a `Dog`.

Now, suppose we have a constructor in the `Dog` class that accepts `name` and `age` as parameters, and a default constructor in the `HomeOwner` class:

```
public class Dog {  
    String name;  
    int age;  
  
    public Dog(String name, int age) {  
        this.name = name;  
        this.age = age;  
    }  
}  
  
public class HomeOwner {  
    String name;  
    Dog pet;  
  
    public HomeOwner() {  
        this.name = "Timothy";  
        this.pet = new Dog("Timothy Jr.", 5);  
    }  
}
```

When we create an instance of `HomeOwner`:

```
HomeOwner owner = new HomeOwner();
```

A new `Dog` object is created and assigned to the `pet` field of the `HomeOwner` object. Importantly, the `pet` field stores the *reference* to the `Dog` object, not the actual `Dog` object data.

This creates a chain of references in memory, with the `owner` variable pointing to the `HomeOwner` object, which in turn has a reference to the `Dog` object.

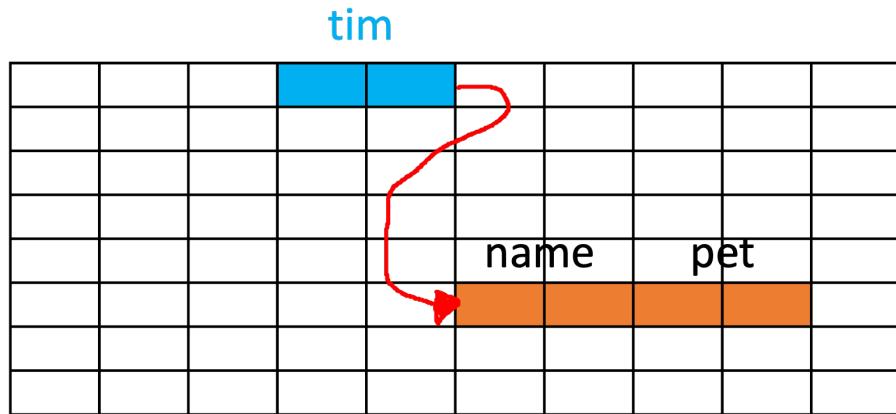


Figure 15.2: HomeOwner Object Creation

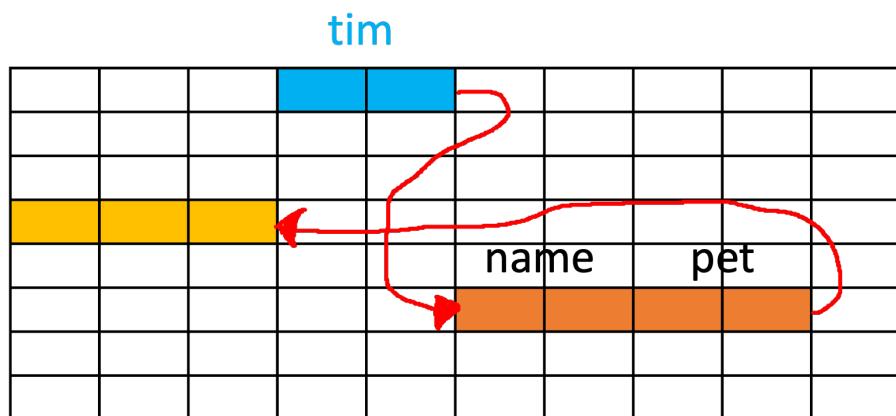


Figure 15.3: Dog Object Creation and Assignment

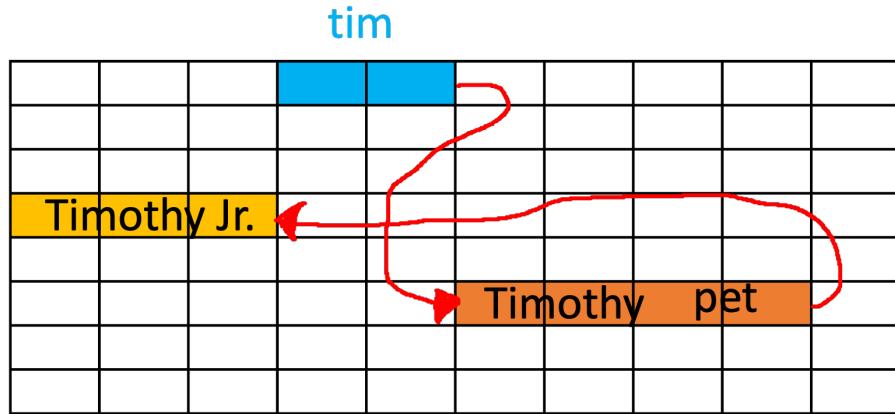


Figure 15.4: Chain of References

This understanding of references and memory allocation sets the foundation for how linked lists work in Java, which we will explore in the upcoming sections.

15.4 Implementing Linked Lists in Java

Absolutely! Last time, we used arrays to construct lists, but we can indeed build a list using the concept of references we have just discussed.

We'll start by creating a class where we'll store a reference to the first piece of data in our List. Then, our first piece of data will contain a reference to the next piece, and the next piece will point to the one following it, and so on.

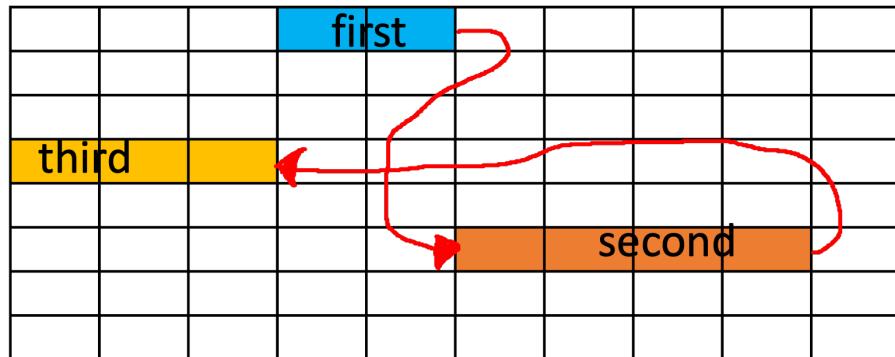


Figure 15.5: Linked List

This gives us the following structure:

```
class LinkedList {
    firstPieceOfData head;
}

class firstPieceOfData {
    nextPieceOfData next;
}
```

However, there's no actual data in this structure yet.

So, how about this?

```
class LinkedList {
    firstPieceOfData head;
}

class firstPieceOfData {
    nextPieceOfData next;
    actualData data;
}

class nextPieceOfData {
    nextPieceOfData next;
    actualData data;
}
```

To ensure our list is flexible and can store any type of data, we should make it generic:

```
class LinkedList<T> {
    firstPieceOfData<T> head;
}

class firstPieceOfData<T> {
    nextPieceOfData<T> next;
    T data;
}

class nextPieceOfData<T> {
    nextPieceOfData<T> next;
    T data;
}
```

Here, `<T>` is a type parameter that allows us to define a class with placeholders for the types

they use, which we can specify when we create an instance of the class.

We can also observe that `firstPieceOfData` and `nextPieceOfData` are structurally the same, as both classes store a reference to an object that holds data, and both hold some actual data. We can therefore simplify our structure by using a single class for both, like so:

```
class LinkedList<T> {
    Node<T> head;
}

class Node<T> {
    Node<T> next;
    T data;
}
```

The list we've created is known as a *singly linked list*. In some algorithms, we may want to access the previous node. In that case, we can store a reference to the previous node as well, creating what's known as a *doubly linked list*:

```
class LinkedList<T> {
    Node<T> head;
    Node<T> tail;
}

class Node<T> {
    Node<T> next;
    Node<T> prev;
    T data;
}
```

When we create a `Node` object, the constructor should initialize the `next` reference to `null` and `data` to the data we pass in. This is because at that point, we don't yet know what the next node will be, but we do know what data we want to store.

```
class SinglyLinkedNode<T> {
    private SinglyLinkedNode<T> next;
    private T data;

    public SinglyLinkedNode(T data) {
        this.data = data;
        this.next = null;
    }
}
```

We also renamed the `Node` class to `SinglyLinkedListNode` to make it clear that this is for a singly-linked list.

The constructor for `SinglyLinkedList` will just set the head to null, initializing it as an empty list.

```
class SinglyLinkedList<T> {
    private SinglyLinkedListNode<T> head;

    public SinglyLinkedList() {
        this.head = null;
    }
}
```

15.5 6.5 Adding to a Singly Linked List

The `SinglyLinkedListNode`

class needs getter and setter methods for accessing and modifying its private data and next fields.

```
class SinglyLinkedListNode<T> {
    private SinglyLinkedListNode<T> next;
    private T data;

    public SinglyLinkedListNode(T data) {
        this.data = data;
        this.next = null;
    }

    public SinglyLinkedListNode<T> getNext() {
        return this.next;
    }

    public void setNext(SinglyLinkedListNode<T> next) {
        this.next = next;
    }

    public T getData() {
        return this.data;
    }
}
```

```
    public void setData(T data) {
        this.data = data;
    }
}
```

With the getters and setters in place, we can now turn our attention to adding a new node to the list. Here's what we need to do when the list is empty:

```
public class SinglyLinkedList<T> {
    private SinglyLinkedNode<T> head;

    public SinglyLinkedList() {
        this.head = null;
    }

    public void add(T data) {
        head = new SinglyLinkedNode<T>(data);
    }
}
```

When we're adding elements or nodes to our singly linked list, we've seen that for the first addition, we simply set the `head` to a new node with the provided data. This is essentially the creation of the first element or node in our list.

Let's take a moment to visualize this scenario:

```
head --> [A]
```

In this ASCII diagram, the `-->` represents the link from the `head` to the first node `A`.

Now, when we're adding a second element, we must be careful. We can't simply overwrite the `head` again, as it contains our first piece of data. Instead, we want to add our new node at the next available position.

We achieve this by setting the `next` member of the `head` node to point to our new data, like so:

```
public void add(T data) {
    if (head == null) {
        head = new SinglyLinkedNode<T>(data);
    } else {
        head.setNext(new SinglyLinkedNode<T>(data));
    }
}
```

```
}
```

The updated state of the list is now:

```
head --> [A] --> [B]
```

But what if we want to add a third node? Using the code above, we would inadvertently overwrite the link from A to B, and instead, link A directly to the new node.

We have to revise our strategy. What if we checked whether the `next` of the head was null and then added our new node there?

```
public void add(T data) {
    if (head == null) {
        head = new SinglyLinkedNode<T>(data);
    } else if (head.getNext() == null) {
        head.setNext(new SinglyLinkedNode<T>(data));
    } else {
        head.getNext().setNext(new SinglyLinkedNode<T>(data));
    }
}
```

This would result in:

```
head --> [A] --> [B] --> [C]
```

That seems to work, right? But this code has limitations; it only accommodates the addition of up to three nodes.

Let's observe what happens when we add a fourth node using this code. We'd end up with:

```
head --> [A] --> [B] --> [D]
```

The third node, [C], disappears! That's not what we want. We need a way to add nodes to our list, regardless of its current length. Can you see a pattern forming in how we're adding nodes? How about a loop to iterate over the list until we find a node with a `null next` field?

```
public void add(T data) {
    if (head == null) {
        head = new SinglyLinkedNode<T>(data);
    } else {
```

```

SinglyLinkedListNode<T> current = head;

while (current.getNext() != null) {
    current = current.getNext();
}

current.setNext(new SinglyLinkedListNode<T>(data));
}
}

```

Now, each time we call `add()`, we start at the head and follow the `next` references until we find a node where `next` is `null`. We then set that node's `next` to the new node, effectively adding it to the end of the list.

And there you have it! We've cracked the code for adding nodes to a singly linked list, regardless of its size.

15.6 Searching a Singly Linked List

Having explored the intricacies of adding elements to our singly linked list, it's only logical to now ask, "How can we find an element in our list?" After all, what good is storing data if we can't retrieve it?

Let's delve into the process of developing a `search` method that behaves similarly to the `indexOf` method in the Java standard library. Our `search` method should return the index of the first occurrence of a given element in the list and `-1` if the element is not found.

Before we get started, there's an important update we need to make to our generics. We'll be comparing objects in our list to the search target, which means we need to ensure these objects are comparable. To do this, we'll update `T` to `T extends Comparable<T>`. This ensures that the type `T` implements the `Comparable` interface, providing us with the ability to compare objects of this type.

```

class SinglyLinkedListNode<T extends Comparable<T>> {
    private SinglyLinkedListNode<T> next;
    private T data;

    //...
}

class SinglyLinkedList<T extends Comparable<T>> {
    private SinglyLinkedListNode<T> head;
}

```

```
//...
}
```

Alright, with this adjustment in place, let's get to developing our `search` method!

If you reflect on how we traverse our linked list, it may become clear that the same process can be utilized for searching. We start at the `head`, and we progress through the list via the `next` pointers until we either find our target or reach the end of the list.

Here's a potential implementation for the `search` method:

```
public int search(T target) {
    SinglyLinkedListNode<T> current = head;
    int index = 0;

    while (current != null) {
        if (current.getData().compareTo(target) == 0) {
            return index;
        }
        index++;
        current = current.getNext();
    }

    return -1; // Target not found
}
```

Let's examine this piece by piece. We start by creating a reference to the `head` of our list and initializing an `index` variable at 0. We then enter a `while` loop that will continue as long as `current` is not `null`, effectively iterating over the entire list.

Inside the loop, we compare the data of the `current` node to our `target` using the `compareTo` method, which is available to us thanks to our `Comparable` constraint. If the data matches our target (`compareTo` returns 0), we've found our target and we return the current index.

If the data does not match our target, we increment our `index` and move to the next node in the list. If we reach the end of the list without finding our target, the method returns `-1`, indicating the target is not in the list.

And that's how we develop a search method for our singly linked list, providing us with the means to locate and retrieve data efficiently!

15.7 Removing from a Singly Linked List

Just as we learned how to add elements to a singly linked list, let's now turn our attention to the process of removing elements. As with adding elements, we'll start with a simpler case and gradually address more complex scenarios. We'll design a `removeLast` method, which removes the last node from the list.

There are three scenarios we need to consider:

1. The list is empty.
2. The list contains only one element.
3. The list contains more than one element.

15.7.1 Case 1: Empty Linked List

The simplest scenario to handle is an empty list. If the list is empty, we have nothing to remove. Here's a simple starting point for our `removeLast` method:

```
public void removeLast() {  
    if (head == null) {  
        return; // Nothing to remove.  
    }  
}
```

This code handles an empty list by simply returning without doing anything.

15.7.2 Case 2: Linked List with One Element

Now, let's consider the case where our list contains only one element.

To remove the only element from the list, we would set `head` to `null`, effectively removing the link to that node.

Here's how our `removeLast` method looks now:

```
public void removeLast() {  
    if (head == null) {  
        return; // Nothing to remove.  
    } else if (head.getNext() == null) {  
        head = null; // Remove the only node in the list.  
    }  
}
```

15.7.3 Case 3: Linked List with More Than One Element

When the list contains more than one element, we must find the last node and the node before it. Why? Because to remove the last node, we must set the `next` field of the node before it to `null`.

Let's think about how we might achieve this. We can't simply set `head.getNext() = null` like we did with the one-element list, because this would leave us with only the first node.

Instead, we could use a loop similar to the one in the `add` method, but with a twist. Instead of stopping when `current.getNext() == null`, which would leave us at the last node, we stop when `current.getNext().getNext() == null`, which will leave us at the second-to-last node.

```
public void removeLast() {
    if (head == null) {
        return; // Nothing to remove.
    } else if (head.getNext() == null) {
        head = null; // Remove the only node in the list.
    } else {
        SinglyLinkedListNode<T> current = head;
        while (current.getNext().getNext() != null) {
            current = current.getNext();
        }
        current.setNext(null); // Remove the last node from the list.
    }
}
```

Now, we have a `removeLast` method that handles any size list!

15.7.4 Moving to `remove(int i)`

Removing the last node from a list is a good starting point, but what if we want to remove an arbitrary node at index `i`?

For this, we would need to iterate over the nodes until we reach the $(i-1)$ th node (just before the node we want to remove), then update its `next` field to skip over the `i`th node and link to the $(i+1)$ th node.

However, this requires careful checking of edge cases, such as when `i` is 0 (requiring us to update the `head`), or when `i` is greater than the size of the list.

This will be our next challenge to tackle! Let's see how we can go about creating a `remove(int i)` method, which will remove a node at a given index from our singly linked list.

As previously mentioned, our approach will be to traverse the list until we reach the $(i-1)$ th node, and then adjust its `next` field to skip the i th node, effectively removing it from the list.

Here's a basic version of this method:

```
public void remove(int i) {
    if (i == 0) {
        head = head.getNext(); // If the node to remove is the head, move the head to the next
    } else {
        SinglyLinkedListNode<T> current = head;
        for (int j = 0; j < i - 1; j++) {
            current = current.getNext();
        }
        current.setNext(current.getNext().getNext());
    }
}
```

This is a simple implementation, but it lacks protection against edge cases. For instance, what happens if i is negative or if it's larger than the size of the list?

We can address these cases by adding a few checks to our method:

```
public void remove(int i) {
    if (i < 0) {
        return; // Do nothing for negative index.
    } else if (i == 0) {
        head = head.getNext(); // If the node to remove is the head, move the head to the next
    } else {
        SinglyLinkedListNode<T> current = head;
        for (int j = 0; j < i - 1; j++) {
            if (current.getNext() == null) {
                return; // If the index is out of range, do nothing.
            }
            current = current.getNext();
        }
        if (current.getNext() != null) {
            current.setNext(current.getNext().getNext());
        }
    }
}
```

This updated method takes care of any negative index or index that's out of range by simply returning without making any changes. If the index is zero, we update the head of the list to

be the next node. If the index is within the range of the list, we find the node at $(i-1)$ and update its `next` reference to skip over the i th node.

And with that, we've covered the basics of how to add and remove nodes from a singly linked list!

15.7.5 Summary and Conclusion

In this chapter, we learned how to implement linked lists in Java using references and objects. We saw how references can point to other objects in memory and how we can use them to create nodes that store data and link to other nodes. We also learned how to define a class for a singly linked list that contains a reference to the first node, called the head, and methods for adding, searching, and removing elements. We compared the advantages and disadvantages of linked lists with array-backed lists and discussed some applications of linked lists in real-world problems. We concluded that linked lists are a dynamic and flexible data structure that can grow and shrink as needed, but they also have some drawbacks such as extra memory overhead, lack of random access, and potential memory leaks.

Part VI

Graphs

16 The Graph Data Structure

16.1 Background and Motivation

Imagine you want to represent the connections between you and your Instagram followers in a data structure. The diagram below (Figure 16.1) shows a simple representation of your followers as numbered vertices and the edges represent the connections between you and your followers.




Figure 16.1: A representation of your Instagram followers.

At first glance, it might appear to be a tree structure, but that is not the case. Your followers can follow other people, who in turn can have their followers. This creates a recursive relationship that cannot be represented using a tree data structure (see Figure 16.2).

To accurately capture this complex relationship, we need to use a graph data structure. Graphs consist of a set of vertices (or nodes) and a set of edges that connect them. In the context of Instagram followers, the vertices represent the users, and the edges represent the connections between them.

Using a graph data structure allows us to represent the recursive nature of the relationship between you and your followers, enabling us to model and analyze the complex network of connections in a more accurate way.

Sure! I will translate the given block according to the specifications you provided. Here's the updated version:

Furthermore, the relationship between you and your followers is even more complex. For instance, you can follow someone who does not follow you back, creating a directed relationship where the edges have a direction. In this case, the edges represent the connections from you to your followers, but not the other way around. Such relationships are often modeled using directed graphs, which are a common use case for graphs. Figure 16.3 visualizes this directed relationship.




Figure 16.2: A representation of your instagram followers where they're allowed to follow other people too.




Figure 16.3: A directed representation of your Instagram followers. Here, an arrow going from vertex A to vertex B indicates that A follows B , but B does not necessarily follow A .

16.2 Introduction

A **graph** is a non-linear data structure that consists of a set of vertices (also called nodes) and a set of edges (or connections) that connect these vertices. In this data structure, the arrangement of vertices and edges allows for a more flexible and complex representation of relationships between data elements compared to linear data structures like arrays, lists, or queues.

The concept of **adjacency** refers to the relationship between two vertices in a graph. If there is an edge connecting two vertices, they are said to be adjacent. **Incidence** is the relationship between a vertex and an edge. A vertex is said to be incident to an edge if it is one of the two vertices connected by that edge.

Graphs have numerous real-life applications, and some examples include:

- Social networks, where vertices represent people and edges represent friendships or connections
- Transportation networks, where vertices represent locations and edges represent roads or routes
- Coronavirus transmission networks, where vertices represent individuals and edges represent transmission paths

16.3 Graph Terminology

Before diving into the implementation of graph data structures, let's discuss some basic terms and properties of graphs.

16.3.1 Basic Terms and Properties

- A **graph** is a data structure for representing connections among items and consists of **vertices** connected by **edges**.
- A **vertex** (or node) represents an item in a graph.
- An **edge** represents a connection between two vertices in a graph.
- Two vertices are **adjacent** if connected by an **edge**.
- **Directed vs Undirected:** In an **undirected graph**, the edges have no specific direction, meaning that if there is an edge between vertices A and B, the connection is mutual. In a **directed graph** (also called a digraph), the edges have a direction, indicating an asymmetrical relationship between vertices. (See Figure 16.4 and Figure 16.5 for examples.)




Figure 16.4: Example of an undirected graph.




Figure 16.5: Example of a directed graph.

- **Weighted vs Unweighted:** In an **unweighted graph**, all edges have equal importance, while in a **weighted graph**, each edge is assigned a value (or weight), representing the importance, cost, or distance between the connected vertices. (See Figure 16.6 and Figure 16.7 for examples.)




Figure 16.6: Example of an unweighted graph.

- **Simple vs Multigraph:** A **simple graph** has no more than one edge between any pair of vertices and does not contain any self-loops (edges that connect a vertex to itself). A **multigraph** can have multiple edges between the same pair of vertices and may include self-loops. (See Figure 16.8 and Figure 16.9 for examples.)
- **Degree:** The degree of a vertex is the number of edges incident to it. In a directed graph, we can distinguish between in-degree (the number of edges directed towards the vertex) and out-degree (the number of edges directed away from the vertex). See Figure 16.10 for an example.
- **Path:** A path in a graph is a sequence of vertices connected by edges. See Figure 16.11 for an example.
- **Cycle:** A cycle is a closed path, where the first and last vertices in the path are the same, and no vertex is visited more than once. See Figure 16.12 for an example.




Figure 16.7: Example of a weighted graph.




Figure 16.8: Example of a simple graph.




Figure 16.9: Example of a multigraph.




Figure 16.10: Example graph with vertex degrees.




Figure 16.11: Example graph with a path from A to D.




Figure 16.12: Example graph with a cycle. (A-B-C)

- **Connected vs Disconnected:** A graph is connected if there is a path between every pair of vertices. If there is at least one pair of vertices with no path between them, the graph is disconnected. See Figure 16.13 and Figure 16.14 for examples.




Figure 16.13: Example of a connected graph.

16.3.2 Graph Notation

We can use a notation like $G(V, E)$, where V is the set of vertices and E is the set of edges, to represent a graph.

16.3.3 Special Types of Graphs

- **Complete Graph:** A complete graph is a simple graph in which every pair of vertices is connected by a unique edge. See Figure 16.15 for an example.
- **Bipartite Graph:** A bipartite graph is a graph whose vertices can be divided into two disjoint sets such that all edges connect vertices from one set to the other, with no edges connecting vertices within the same set. See Figure 16.16 for an example.




Figure 16.14: Example of a disconnected graph.




Figure 16.15: Example of a complete graph.




Figure 16.16: Example of a bipartite graph.

- **Tree:** A tree is an undirected graph with no cycles, and all vertices are connected. It has a hierarchical structure, with one vertex acting as the root, and the other vertices connected in a parent-child relationship. See Figure 16.17 for an example.




Figure 16.17: Example of a tree.

16.4 Graph Representation

In order to work with graphs in code or store them in memory, we need efficient ways to represent them. There are multiple methods to represent graphs, and the choice of representation depends on factors such as the density of the graph, the operations to be performed, and memory constraints.

In this section, we will discuss two common methods to represent a graph: adjacency list and adjacency matrix.

16.4.1 Adjacency List

An **adjacency list** represents a graph by storing a list of adjacent vertices for each vertex in the graph. This can be implemented using an array of lists or a hash table, where the index

or key corresponds to a vertex, and the value is a list of adjacent vertices.




Figure 16.18: See adjacency list for this example Listing 16.1

Adjacency list representation for figure Figure 16.18:

The adjacency list representation is efficient for sparse graphs (graphs with relatively few edges) as it only stores the existing edges, reducing memory usage. This representation also allows for faster traversal of a vertex's neighbors.


16.4.2 Adjacency Matrix

An **adjacency matrix** is a two-dimensional array (or matrix) where the cell at the i -th row and j -th column represents the edge between vertex i and vertex j . For an undirected graph, the adjacency matrix is symmetric. For a weighted graph, the values in the cells represent the weights of the edges; for an unweighted graph, the cells contain either 1 (edge exists) or 0 (no edge).

Listing 16.1 Adjacency list representation.

```
A: [B, C]
B: [A, D]
C: [A, D]
D: [B, C]

// or as an arraylist of arraylists -
[[B, C], [A, D], [A, D], [B, C]]
// here, the index of the outer arraylist represents the vertex.
// in order for this to work, the order of the vertices must be
// fixed, and stored separately.
```



See adjacency matrix for this example Listing 16.2

Adjacency matrix representation (unweighted):

The adjacency matrix representation is suitable for dense graphs (graphs with many edges) or when checking for the presence of an edge between two vertices needs to be fast. However, this representation can be inefficient in terms of memory usage, especially for large, sparse graphs, as it stores information for all possible edges, even if they do not exist.

Listing 16.2 Adjacency matrix representation.

```
A B C D
A 0 1 1 0
B 1 0 0 1
C 1 0 0 1
D 0 1 1 0
```

16.4.3 Converting Between Representations

To convert a graph diagram or notation into an adjacency list or an adjacency matrix, follow these steps:

1. Identify the vertices and edges in the graph.
2. For an adjacency list, create an empty list or hash table for each vertex. For each edge, add the adjacent vertices to the corresponding lists.
3. For an adjacency matrix, create a square matrix with dimensions equal to the number of vertices. For each edge, set the corresponding cells in the matrix to 1 (or the edge weight for weighted graphs).

To convert an adjacency list or an adjacency matrix back into a graph diagram or notation, follow these steps:

1. Identify the vertices based on the keys (for an adjacency list) or indices (for an adjacency matrix).
2. For an adjacency list, iterate through the lists and draw an edge for each adjacent vertex.
3. For an adjacency matrix, iterate through the matrix cells and draw an edge for each non-zero value (or the corresponding weight for weighted graphs).

16.5 Graph Traversal

Imagine you want to find the average age of all users on Facebook. With billions of users, it is infeasible to hold the entire graph of the friend network in memory. Ideally, we would want to find out information on each user one at a time, on a per-need basis. To achieve this, we can use graph traversal algorithms, which allow us to visit each user, add up their ages, and then calculate the average. A simple way to do this is to load information on a user, add all their friends to a stack, and then keep popping from the stack and requesting data from Facebook for each friend. When we receive the data, we mark that user as visited to avoid recounting their age if we reach the same user again. We then add friends of each loaded user to our stack and keep repeating until we run out of users in our stack.

This problem illustrates the importance of graph traversal, a fundamental operation in graph theory. Graphs are a powerful and versatile data structure that can model various kinds of relationships and networks, such as social networks, computer networks, transportation networks, web pages, games, and many other domains. Graph traversal allows us to explore and manipulate graphs in various ways, with applications in domains like searching for specific nodes, finding the shortest path between nodes, and analyzing the structure of a graph.

Graph traversal algorithms typically begin with a start node and attempt to visit the remaining nodes from there. They must deal with several troublesome cases, such as unreachable nodes, revisited nodes, and choosing which node to visit next among several options. To handle these cases, graph traversal algorithms use different strategies and data structures to keep track of which nodes have been visited and which nodes are still pending. The most common graph traversal algorithms are breadth-first search (BFS) and depth-first search (DFS), which differ in the order in which they visit the nodes.

In some situations, we may not know the entire graph at once and instead only have access to a node object and its adjacent nodes. As demonstrated in the Facebook example, graph traversal algorithms can be used to solve problems that involve large and dynamic graphs by visiting each user and analyzing their information on a per-need basis.

There are two common methods to traverse a graph:

1. Breadth-First Search (BFS)
2. Depth-First Search (DFS)

By understanding and implementing these graph traversal methods, you can efficiently explore and manipulate complex graphs to solve a wide range of problems.

16.5.1 Breadth-First Search (BFS)

Breadth-First Search explores a graph by visiting all the neighbors of the starting vertex before moving on to their neighbors. BFS uses a queue data structure to keep track of the vertices to visit.

Here's a step-by-step example of BFS traversal (for the graph in example Figure 16.19):

BFS traversal starting from vertex A:

1. Visit A and add its neighbors B and C to the queue: [B, C]
2. Visit B and add its unvisited neighbor D to the queue: [C, D]
3. Visit C and add its unvisited neighbor E to the queue: [D, E]
4. Visit D: [E]
5. Visit E: []




Figure 16.19: Example graph for BFS traversal.

BFS traversal order: A, B, C, D, E

BFS pseudocode:

```
BFS(graph, start):
    Initialize an empty queue Q
    Mark start as visited
    Enqueue start into Q

    while Q is not empty:
        vertex = Dequeue(Q)
        Visit vertex

        for each neighbor of vertex:
            if neighbor is not visited:
                Mark neighbor as visited
                Enqueue neighbor into Q
```

16.5.2 Depth-First Search (DFS)

Depth-First Search explores a graph by visiting a vertex and its neighbors as deeply as possible before backtracking. DFS can be implemented using recursion or an explicit stack data structure.

Here's a step-by-step example of DFS traversal (for the graph in example Figure 16.20):

DFS traversal starting from vertex A:

1. Visit A and recurse on its first neighbor B
2. Visit B and recurse on its first neighbor D
3. Visit D and backtrack (no unvisited neighbors)
4. Backtrack to A and recurse on its next neighbor C
5. Visit C and recurse on its first neighbor E
6. Visit E and backtrack (no unvisited neighbors)

DFS traversal order: A, B, D, C, E

DFS pseudocode (recursive):

```
DFS(graph, vertex):
    Mark vertex as visited
    Visit vertex

    for each neighbor of vertex:
```




Figure 16.20: Example graph for DFS traversal.

```

if neighbor is not visited:
    DFS(graph, neighbor)

```

DFS pseudocode (iterative with a stack):

```

DFS(graph, start):
    Initialize an empty stack S
    Mark start as visited
    Push start onto S

    while S is not empty:
        vertex = Pop(S)
        Visit vertex

        for each neighbor of vertex:
            if neighbor is not visited:
                Mark neighbor as visited
                Push neighbor onto S

```

16.5.3 Applications and Variations of BFS and DFS

Both BFS and DFS have numerous applications and can be adapted to solve various graph-related problems:

- **Shortest path:** BFS can be used to find the shortest path between two vertices in an unweighted graph. The algorithm can be modified to keep track of the path length or the actual path itself.
- **Connected components:** Both BFS and DFS can be used to find connected components in an undirected graph. By running the traversal algorithm and marking visited vertices, we can identify the set of vertices reachable from a starting vertex. Repeating this process for all unvisited vertices will find all connected components in the graph.
- **Topological sorting:** DFS can be adapted to perform a topological sort on a directed acyclic graph (DAG). A topological ordering is a linear ordering of the vertices such that for every directed edge (u, v) , vertex u comes before vertex v in the ordering. This can be useful in scheduling tasks with dependencies or determining the order of courses in a curriculum.
- **Bipartite graph check:** BFS or DFS can be used to check if a graph is bipartite. The algorithm can be modified to color vertices while traversing the graph. If at any point during the traversal, two adjacent vertices have the same color, the graph is not bipartite.

- **Graph cycle detection:** DFS can be used to detect cycles in a graph. By keeping track of the recursion stack, we can determine if a vertex is visited more than once in the same path, indicating a cycle.

In summary, graph traversal is a fundamental operation in graph theory with various applications. Breadth-First Search (BFS) and Depth-First Search (DFS) are two common techniques to traverse a graph, each with its own advantages and use cases. Understanding these algorithms and their variations can help solve a wide range of graph-related problems.

Part VII

Hashing, HashMaps and HashTables

17 Hashing, Hash Tables, and Hash Maps

17.1 Background and Motivation

17.1.1 Indexing

Indexing refers to the idea of accessing a certain element of an array by referring to it using a specific number, called the index. Using the same index always returns the same element, as long as the array remains unchanged. For example, to access the 12th element of an array `arr`, we use `arr[11]`.

The math for finding the address of an element in the array works out to be:

```
baseAddress + (index * sizeOfElement)
```

Now, keeping that in mind, let's explore the limitations of indexing.

17.1.2 Limitations of Indexing

Suppose we want to access an element in the array using a string as an index, such as `arr["dhruv"]`. What is stopping us?

The problem is that we cannot calculate `baseAddress + ("dhruv" * sizeOfElement)` because the index, in this case, is a string, not a number. The operation is not defined, and therefore, we can't directly use strings as indices in an array.

17.1.3 Mapping Strings to Numbers

Let's consider an array of strings. We can use it to map a number to a string:

```
0 -> "Alice"  
1 -> "Bob"  
2 -> "Charlie"
```

If we can use an array to map a number to a string, can we also use it to map strings to numbers? Yes, we can!

One way to do this is by searching linearly for a string in the array to find its index. For example, we can find the string “Dhruv” at index 5:

```
0 -> "Alice"  
1 -> "Bob"  
2 -> "Charlie"  
...  
5 -> "Dhruv"
```

The index at which we found the string “Dhruv” (in this case, 5) can be used as a key in a different array to find the data related to “Dhruv”. However, this method of linear searching can be quite slow for large datasets.

17.1.4 Hashing: A Better Solution

This is where **hashing** comes into play. Hashing allows us to efficiently map strings (or any other non-numeric keys) to numbers. By using a hash function, we can convert a string into a number that represents the index in the array.

A **hash function** takes a key as input and outputs an index in the hash table’s array. A good hash function has the following criteria:

- **Uniform distribution:** The hash function should distribute keys evenly across the array to minimize collisions (when multiple keys map to the same index).
- **Minimal collisions:** A good hash function should minimize the chance of collisions.
- **Fast computation:** The hash function should be fast to compute, allowing for quick insertion, deletion, and retrieval of data.
- **Deterministic output:** The hash function should produce the same output for the same input every time it is called.

For example, let’s consider a simple hash function that converts the first character of a string into its ASCII code:

```
hash("dhruv") = ASCII('d') = 100
```

The output of the hash function is 100, which we can use as an index in an array to store or retrieve data related to “dhruv”. This allows us to use strings (and other non-numeric keys) as indices, achieving our goal of efficient mapping.

17.2 Hash Functions

17.2.1 Introduction

Previously, we managed to map a string “D” to some data, just like an index maps to some data in an array. The resulting data structure that can map any string to any data is called a **hash table**. The function used to map strings to data is called a **hash function**. The concept of mapping “D” to some data can be referred to as **hashing** “D” to an index 3, which is where we found the value corresponding to “D”.

Now, we’ll learn about a way of mapping any object (called a key) to any other object (called the record, or the value). For instance, your student ID can be a key, and all the data about you on your ID can be stored in a record object.

17.2.2 Hashing

Hashing can be thought of as a method for storing and retrieving records from a database. It lets you insert, delete, and search for records based on a search key value. When properly implemented, these operations can be performed in constant time. In fact, a properly tuned hash system typically looks at only one or two records for each search, insert, or delete operation. This is far better than the $O(\log n)$ average cost required to do a binary search on a sorted array of n records, or the $O(\log n)$ average cost required to do an operation on a binary search tree. However, even though hashing is based on a very simple idea, it is surprisingly difficult to implement properly. Designers need to pay careful attention to all of the details involved with implementing a hash system.

A hash system stores records in an array called a **hash table**, which we will call HT . Hashing works by performing a computation on a search key K in a way that is intended to identify the position in HT that contains the record with key K . The function that does this calculation is called the **hash function**, and will be denoted by the letter h . Since hashing schemes place records in the table in whatever order satisfies the needs of the address calculation, records are not ordered by value. A position in the hash table is also known as a **slot**. The number of slots in hash table HT will be denoted by the variable M with slots numbered from 0 to $M-1$. The goal for a hashing system is to arrange things such that, for any key value K and some hash function h , $i = h(K)$ is a slot in the table such that $0 \leq i < M$, and we have the key of the record stored at $HT[i]$ equal to K .

Hashing is not good for applications where multiple records with the same key value are permitted. Hashing is not a good method for answering range searches. In other words, we cannot easily find all records (if any) whose key values fall within a certain range. Nor can we easily find the record with the minimum or maximum key value or visit the records in key order. Hashing is most appropriate for answering the question, ‘What record, if any, has

key value K?’ For applications where all search is done by exact-match queries, hashing is the search method of choice because it is extremely efficient when implemented correctly.

Hashing generally takes records whose key values come from a large range and stores those records in a table with a relatively small number of slots. Since keys have a large range and values have smaller, limited slots for storage – A hash function might sometimes end up hashing two keys to the same slot. We refer to such an event as a **collision**.

To illustrate, consider a classroom full of students. What is the probability that some pair of students shares the same birthday (i.e., the same day of the year, not necessarily the same year)? If there are 23 students, then it is unlikely that more than one student will share the same birthday. There are 365 “slots” or possible days a student can have a birthday on; but only 23 “keys”. As the number of students increases, the probability of a “collision” or two students sharing a birthday increases. To be practical, a database organized by hashing must store records in a hash table that is not so large that it wastes space.

We would like to pick a hash function that maps keys to slots in a way that makes each slot in the hash table have equal probability of being filled for the actual set keys being used. Unfortunately, we normally have no control over the distribution of key values for the actual records in a given database or collection. So how well any particular hash function does depends on the actual distribution of the keys used within the allowable key range. In some cases, incoming data are well distributed across their key range. For example, if the input is a set of random numbers selected uniformly from the key range, any hash function that assigns the key range so that each slot in the hash table receives an equal share of the range will likely also distribute the input records uniformly within the table.

However, in many applications the incoming records are highly clustered or otherwise poorly distributed. When input records are not well distributed throughout the key range it can be difficult to devise a hash function that does a good job of distributing the records throughout the table, especially if the input distribution is not known in advance. For example, If the input is a collection of English words, the beginning letter will be poorly distributed. A dictionary of words mapped to their frequency is often used in rudimentary natural language processing algorithms.

In conclusion, anything can be a hash function (i.e., map a value to an index), but not everything can be a good hash function. A function that always returns the index 0 is a hash function that maps everything to 0. It’s no good but it’s still a hash function. An example of a commonly used hash function is the modulus operator! It is common for N-sized hash tables to use the modulus of N as a hash function. If N is 20, data for 113 will be hashed to index $113 \% 20 = 13$.

But if we use the modulo operator as a hash function, what do we do when multiple pieces of data map to the same index? $53 \% 20 = 13$, $73 \% 20 = 13$, etc. But if you think about it, we can store everything at 13! By using nested data structures... More on this later.

17.2.3 Simple Hash Functions

Let's apply a simple hash function to a set of keys and compute their indices. In this example, we'll use the modulo operation as the hash function. Given a hash table with a size of 5, we can compute the indices for the keys as follows:

```
HashTable size: 5
HashFunction: key % size

Keys: 15, 28, 47, 10, 33

Indices:
15 % 5 = 0
28 % 5 = 3
47 % 5 = 2
10 % 5 = 0
33 % 5 = 3
```

17.2.4 Other Types of Hash Functions

17.2.4.1 Direct Hashing

A direct hash function uses the item's key as the bucket index. For example, if the key is 937, the index is 937. A hash table with a direct hash function is called a direct access table. Given a key, a direct access table search algorithm returns the item at index key if the bucket is not empty, and returns null (indicating item not found) if empty.

Limitations:

A direct access table has the advantage of no collisions: Each key is unique (by definition of a key), and each gets a unique bucket, so no collisions can occur. However, a direct access table has two main limitations:

1. All keys must be non-negative integers, but for some applications, keys may be negative.
2. The hash table's size equals the largest key value plus 1, which may be very large.

Similarly, there are other hash functions each with their own characteristics.

17.2.4.2 Modulo Hash

A modulo hash function computes the index by taking the remainder of the key divided by the table size M. This is a simple and effective way to convert a large key range into a smaller index range. The hash function can be defined as:

$$h(K) = K \% M$$

17.2.4.3 Mid-Square Hash

A mid-square hash function computes the index by first squaring the key, and then extracting a portion of the squared value as the index. This approach is especially useful when the keys are not uniformly distributed. The hash function can be defined as:

$$h(K) = \text{middle_digits}(K^2)$$

17.2.4.4 Mid-Square Hash with Base 2

A mid-square hash function with base 2 is a variation of the mid-square hash function, where the key is first squared, and then the middle bits of the binary representation of the squared value are extracted as the index. This approach is especially useful for binary keys. The hash function can be defined as:

$$h(K) = \text{middle_bits}(K^2)$$

17.2.4.5 Multiplicative String Hashing

A multiplicative string hashing function computes the index by treating the characters in the string as numbers and combining them using a multiplication and a constant. This approach can help achieve a good distribution of string keys in the hash table. The hash function can be defined as:

$$h(K) = (c_1 * a^{(n-1)} + c_2 * a^{(n-2)} + \dots + c_n) \% M$$

where c_1, c_2, \dots, c_n are the character codes of the string, a is a constant, n is the length of the string, and M is the size of the hash table.

Here's the ASCII representation of the resulting hash table:

| Index | | Key |
|-------|--|-----|
| 0 | | 15 |
| 1 | | - |
| 2 | | 47 |
| 3 | | 28 |
| 4 | | - |

In this example, we can see that the keys 15 and 10, as well as 28 and 33, have collided, as they both map to the same indices (0 and 3, respectively).

17.2.5 Trade-offs Between Different Hash Functions

There are trade-offs between different hash functions in terms of performance and complexity:

- A simple hash function, like the modulo operation, is fast to compute but may not distribute keys uniformly, leading to more collisions and reduced performance.
- More complex hash functions, such as cryptographic hash functions, can provide a better distribution of keys but may be slower to compute.

In practice, the choice of a hash function depends on the specific requirements of the application and the data being stored. The goal is to find a balance between uniform distribution, minimal collisions, fast computation, and deterministic output.

17.3 Hash Collisions

Hash collisions occur when two or more keys map to the same index in the hash table. Due to the pigeonhole principle, hash collisions are inevitable, as there are typically more possible keys than available indices in the array. Collisions negatively impact the efficiency of hashing, as they can lead to longer access times for insertion, deletion, and retrieval of key-value pairs.

There are two primary methods to resolve hash collisions: **chaining** and **open addressing**.

17.4 Chaining

Chaining is a collision resolution technique that uses a linked list or another data structure to store multiple key-value pairs at the same index. When a collision occurs, the new key-value pair is simply added to the data structure at the index.

17.4.1 Insertion, Search, and Deletion

Here's how to perform insertion, search, and deletion operations using chaining:

1. **Insertion:** Calculate the index using the hash function. If the index is empty, create a new data structure (e.g., linked list) and insert the key-value pair. If the index is not empty, add the key-value pair to the existing data structure.
2. **Search:** Calculate the index using the hash function. If the index is empty, the key is not in the hash table. If the index is not empty, search the data structure at the index for the key.
3. **Deletion:** Calculate the index using the hash function. If the index is empty, the key is not in the hash table. If the index is not empty, search the data structure at the index for the key and remove it if found.

17.4.2 Advantages and Disadvantages of Chaining

Chaining has several advantages and disadvantages:

- **Advantages:**
 - Easy implementation: Chaining can be easily implemented using existing data structures like linked lists.
 - Dynamic size: The data structure at each index can grow or shrink as needed, allowing for efficient use of space.
- **Disadvantages:**
 - Extra space: Chaining requires additional space for the data structure at each index, which can increase memory overhead.
 - Variable access time: The access time for key-value pairs depends on the length of the data structure at the index, which can vary.

Chaining is a popular method for resolving hash collisions due to its simplicity and dynamic size. However, it may not be the most efficient option for all use cases, especially when memory overhead and variable access times are critical factors.

17.5 Open Addressing

Open addressing is a collision resolution technique that finds an alternative index for a key-value pair if the original index is occupied. When a collision occurs, the algorithm searches for the next available index using a probing technique. There are three common types of probing techniques: linear probing, quadratic probing, and double hashing.

17.5.1 Probing Techniques

1. **Linear probing:** When a collision occurs, search the hash table linearly (one index at a time) until an empty slot is found.
2. **Quadratic probing:** When a collision occurs, search the hash table quadratically (by increasing the index by the square of the probe number) until an empty slot is found.
3. **Double hashing:** When a collision occurs, use a secondary hash function to compute a new index for the key-value pair, and repeat this process until an empty slot is found.

17.5.2 Insertion, Search, and Deletion

Here's how to perform insertion, search, and deletion operations using open addressing:

1. **Insertion:** Calculate the index using the hash function. If the index is empty, insert the key-value pair. If the index is occupied, use the chosen probing technique to find the next available index and insert the key-value pair there.
2. **Search:** Calculate the index using the hash function. If the index is empty, the key is not in the hash table. If the index is occupied, check if the key matches the stored key. If not, use the chosen probing technique to search for the next index until the key is found or an empty index is encountered.
3. **Deletion:** Calculate the index using the hash function. If the index is empty, the key is not in the hash table. If the index is occupied and the key matches the stored key, remove the key-value pair and mark the index as deleted. Continue searching using the chosen probing technique to handle cases where the removed key-value pair was part of a cluster.

17.5.3 Advantages and Disadvantages of Open Addressing

Open addressing has several advantages and disadvantages:

- **Advantages:**
 - No extra space: Open addressing does not require additional space for data structures at each index, making it more memory-efficient.
 - Fixed size: The hash table has a fixed size, which can be useful when memory is limited.
- **Disadvantages:**
 - Clustering: Probing techniques can cause clusters of key-value pairs to form, leading to increased access times.

- Deletion issues: Deleting key-value pairs can create complications, as it may leave “holes” in clusters that need to be addressed.

Open addressing is an alternative method for resolving hash collisions that can be more memory-efficient than chaining. However, it may not be the best option for all use cases, especially when clustering and deletion issues are critical factors.

17.6 Complexity and Load Factor

When analyzing the complexity of hash functions and hash tables, we need to consider the time taken for searching, inserting, or deleting an element. There are two main steps involved in these operations:

1. Computing the hash function for the given key.
2. Traversing the list of key-value pairs present at the computed index.

17.6.1 Time Complexity of Hash Computation

For the first step, the time taken depends on the key and the hash function. For example, if the key is a string “abcd”, then its hash function may depend on the length of the string. But for very large values of n , the number of entries into the map, the length of the keys is almost negligible in comparison to n , so hash computation can be considered to take place in constant time, i.e., $O(1)$.

17.6.2 Time Complexity of List Traversal

For the second step, traversal of the list of key-value pairs present at that index needs to be done. In the worst case, all the n entries are at the same index, resulting in a time complexity of $O(n)$. However, enough research has been done to make hash functions uniformly distribute the keys in the array, so this almost never happens.

17.6.3 Load Factor

On average, if there are n entries and b is the size of the array, there would be n/b entries at each index. This value n/b is called the load factor, which represents the load on our map. The load factor is denoted by the symbol :

$$= n/b$$

This load factor needs to be kept low so that the number of entries at one index is less, and the complexity remains almost constant, i.e., $O(1)$.

17.6.4 Balancing Load Factor and Complexity

To maintain the load factor at an acceptable level, the hash table can be resized when the load factor exceeds a certain threshold. This helps to keep the complexity of hash table operations near $O(1)$ by redistributing the keys uniformly across a larger array.

In conclusion, understanding the complexity and load factor of hash functions is crucial for designing efficient hash tables. By carefully choosing a suitable hash function and managing the load factor, it's possible to achieve near-constant time complexity for various hash table operations.

17.7 Rehashing

Rehashing, as the name suggests, means hashing again. When the load factor increases to more than its pre-defined value (the default value of the load factor is 0.75), the complexity increases. To overcome this issue, the size of the array is increased (typically doubled) and all the values are hashed again and stored in the new, larger array. This helps maintain a low load factor and low complexity.

17.7.1 Why?

Rehashing is done because whenever key-value pairs are inserted into the map, the load factor increases, which implies that the time complexity also increases, as explained earlier. This might not provide the desired time complexity of $O(1)$. Hence, rehashing must be performed, increasing the size of the `bucketArray` to reduce the load factor and the time complexity.

17.7.2 How?

Rehashing can be done as follows:

1. For each addition of a new entry to the map, check the load factor.
2. If the load factor is greater than its pre-defined value (or the default value of 0.75 if not given), then perform rehashing.
3. To rehash, create a new array of double the previous size and make it the new `bucketArray`.
4. Traverse each element in the old `bucketArray` and call the `insert()` method for each, to insert it into the new larger `bucketArray`.

The following diagram illustrates the rehashing process:

Initial bucketArray (size = 4):

| | | | | | | |
|---|----|----|---|---|---|---|
| + | - | - | - | - | - | - |
| | K1 | K2 | | | | |
| + | - | - | - | - | - | - |

After inserting a new key K3 (load factor > 0.75):

New bucketArray (size = 8):

| | | | | | | | | |
|---|----|----|---|----|---|---|---|---|
| + | - | - | - | - | - | - | - | - |
| | K1 | K2 | | K3 | | | | |
| + | - | - | - | - | - | - | - | - |

By rehashing, the hash table maintains its desired time complexity of O(1) even as the number of elements increases. It is important to note that rehashing can be a costly operation, especially if the number of elements in the hash table is large. However, since rehashing is done infrequently and only when the load factor surpasses a certain threshold, the amortized cost of rehashing remains low, allowing the hash table operations to maintain near-constant time complexity.

17.8 Hash Tables vs Hash Maps

Hash tables and **hash maps** differ in their implementation and functionality.

- **Hash tables** use direct hashing, where the key is an integer or can be directly converted to an integer (e.g., a string of digits). The integer is then used to compute the index in the hash table.
- **Hash maps** use indirect hashing, where the key can be any data type. A separate hash function is needed to convert the key into an index in the hash table.

When deciding whether to use a hash table or a hash map, consider the problem domain and the data type of the keys:

- If the keys are integers or can be directly converted to integers, a **hash table** may be a more suitable choice. For example, if you're working with student IDs as keys, a hash table would be a good fit.
- If the keys are of any other data type or cannot be directly converted to integers, a **hash map** would be more appropriate. For example, if you're working with strings, such as usernames or URLs, a hash map would be a better choice.

17.9 HashMaps in Java

A **HashMap** is a collection in Java that implements the Map interface and uses a hash table for storage. It stores key-value pairs, where each key is unique, and the keys are not ordered.

Here's how to use a HashMap in Java:

1. **Import the HashMap class:** To use the HashMap class in your Java code, you'll need to import it from the `java.util` package:

```
import java.util.HashMap;
```

2. **Create a HashMap:** To create a new HashMap, use the following syntax:

```
HashMap<String, Integer> myMap = new HashMap<String, Integer>();
```

3. **Add elements:** To add key-value pairs to the HashMap, use the `put()` method:

```
myMap.put("apple", 3);
myMap.put("banana", 5);
myMap.put("orange", 2);
```

4. **Access elements:** To access the value associated with a key, use the `get()` method:

```
int apples = myMap.get("apple"); // 3
int oranges = myMap.get("orange"); // 2
```

5. **Remove elements:** To remove a key-value pair from the HashMap, use the `remove()` method:

```
myMap.remove("banana");
```

6. **Check if a key exists:** To check if a key is in the HashMap, use the `containsKey()` method:

```
boolean hasApple = myMap.containsKey("apple"); // true
boolean hasGrape = myMap.containsKey("grape"); // false
```

7. **Iterate over keys:** To iterate over the keys in a HashMap, you can use a for-each loop with the `keySet()` method:

```
for (String fruit : myMap.keySet()) {  
    System.out.println(fruit + ": " + myMap.get(fruit));  
}
```

8. **Iterate over values:** To iterate over the values in a HashMap, you can use a for-each loop with the `values()` method:

```
for (Integer count : myMap.values()) {  
    System.out.println(count);  
}
```

9. **Iterate over key-value pairs:** To iterate over the key-value pairs in a HashMap, you can use a for-each loop with the `entrySet()` method:

```
for (HashMap.Entry<String, Integer> entry : myMap.entrySet()) {  
    System.out.println(entry.getKey() + ": " + entry.getValue());  
}
```

A HashMap can be a useful data structure when you need to store key-value pairs efficiently. It provides constant-time performance for common operations like put, get, and remove, making it an ideal choice for various applications.

17.10 HashTables in Java

A **HashTable** is a collection in Java that implements the Map interface and uses a hash table for storage. It is similar to a HashMap but with some differences, such as being synchronized, which makes it thread-safe. HashTable stores key-value pairs, where each key is unique, and the keys are not ordered.

Here's how to use a HashTable in Java:

1. **Import the HashTable class:** To use the HashTable class in your Java code, you'll need to import it from the `java.util` package:

```
import java.util.Hashtable;
```

2. **Create a HashTable:** To create a new HashTable, use the following syntax:

```
Hashtable<String, Integer> myTable = new Hashtable<String, Integer>();
```

3. **Add elements:** To add key-value pairs to the HashTable, use the `put()` method:

```
myTable.put("apple", 3);
myTable.put("banana", 5);
myTable.put("orange", 2);
```

4. **Access elements:** To access the value associated with a key, use the `get()` method:

```
int apples = myTable.get("apple"); // 3
int oranges = myTable.get("orange"); // 2
```

5. **Remove elements:** To remove a key-value pair from the HashTable, use the `remove()` method:

```
myTable.remove("banana");
```

6. **Check if a key exists:** To check if a key is in the HashTable, use the `containsKey()` method:

```
boolean hasApple = myTable.containsKey("apple"); // true
boolean hasGrape = myTable.containsKey("grape"); // false
```

7. **Iterate over keys:** To iterate over the keys in a HashTable, you can use a for-each loop with the `keySet()` method:

```
for (String fruit : myTable.keySet()) {  
    System.out.println(fruit + ": " + myTable.get(fruit));  
}
```

8. **Iterate over values:** To iterate over the values in a HashTable, you can use a for-each loop with the `values()` method:

```
for (Integer count : myTable.values()) {  
    System.out.println(count);  
}
```

9. **Iterate over key-value pairs:** To iterate over the key-value pairs in a HashTable, you can use a for-each loop with the `entrySet()` method:

```
for (Hashtable.Entry<String, Integer> entry : myTable.entrySet()) {  
    System.out.println(entry.getKey() + ": " + entry.getValue());  
}
```

A HashTable can be a useful data structure when you need to store key-value pairs and require thread-safe operations. However, it has some performance overhead due to synchronization, so if thread safety is not a concern, a HashMap is generally a more efficient choice.

17.11 HashSets in Java

A **HashSet** is a collection in Java that implements the Set interface and uses a hash table for storage. It does not store key-value pairs like hash tables or hash maps, but instead stores unique elements. The elements in a HashSet are not ordered, and duplicate values are not allowed.

Here's how to use a HashSet in Java:

1. **Import the HashSet class:** To use the HashSet class in your Java code, you'll need to import it from the `java.util` package:

```
import java.util.HashSet;
```

2. **Create a HashSet:** To create a new HashSet, use the following syntax:

```
HashSet<String> mySet = new HashSet<String>();
```

3. **Add elements:** To add elements to the HashSet, use the `add()` method:

```
mySet.add("apple");
mySet.add("banana");
mySet.add("orange");
```

4. **Remove elements:** To remove elements from the HashSet, use the `remove()` method:

```
mySet.remove("banana");
```

5. **Check if an element exists:** To check if an element is in the HashSet, use the `contains()` method:

```
boolean hasApple = mySet.contains("apple"); // true
boolean hasGrape = mySet.contains("grape"); // false
```

6. **Iterate over elements:** To iterate over the elements in a HashSet, you can use a for-each loop:

```
for (String fruit : mySet) {
    System.out.println(fruit);
}
```

A HashSet can be a useful data structure when you need to store a collection of unique elements without any specific order. It provides constant-time performance for common operations like add, remove, and contains, making it an efficient choice for many applications.

17.12 hashCode and equals in Java

In Java, the `hashCode` method is part of the `Object` class, which is the superclass of all Java classes. The purpose of the `hashCode` method is to provide a default implementation for generating hash codes, which are integer values that represent the memory address of an object.

17.12.1 The hashCode Method

The `hashCode` method has the following signature:

```
public int hashCode()
```

This method returns an integer hash code for the object on which it is called. By default, it returns a hash code that is based on the object's memory address, but this behavior can be overridden in subclasses to provide custom hash code generation.

A well-implemented `hashCode` method should follow these general rules:

1. If two objects are equal according to their `equals()` method, they must have the same hash code.
2. If two objects have the same hash code, they are not necessarily equal according to their `equals()` method.
3. The hash code of an object should not change over time unless the information used in the `equals()` method also changes.

17.12.2 Overriding the hashCode Method

When creating custom classes, it is important to override the `hashCode` method if the `equals()` method is also overridden. This ensures that the general contract of the `hashCode` method is maintained, which is essential for the correct functioning of hash-based data structures like `HashSet` and `HashMap`.

Here's an example of a custom `Person` class that overrides both the `equals()` and `hashCode()` methods:

```
public class Person {  
    private String name;  
    private int age;  
  
    // Constructor, getters, and setters
```

```

@Override
public boolean equals(Object obj) {
    if (this == obj) {
        return true;
    }
    if (obj == null || getClass() != obj.getClass()) {
        return false;
    }
    Person person = (Person) obj;
    return age == person.age && Objects.equals(name, person.name);
}

@Override
public int hashCode() {
    return Objects.hash(name, age);
}
}

```

In this example, the `equals()` method checks if two `Person` objects have the same name and age. The `hashCode()` method uses the `Objects.hash()` utility method, which generates a hash code based on the name and age fields.

17.12.3 Using hashCode with Java Collections

The `hashCode` method plays a crucial role in the performance of Java's hash-based data structures, such as `HashSet`, `HashMap`, and `HashTable`. These data structures rely on the `hashCode` method to efficiently store and retrieve objects based on their hash codes.

When working with these collections, it is important to ensure that the `hashCode` method is correctly implemented for the objects being stored. Failing to do so can lead to poor performance or incorrect behavior.

In summary, the `hashCode` method in Java is a critical part of the `Object` class that provides a default implementation for generating hash codes. When creating custom classes, it is essential to override the `hashCode` method if the `equals()` method is also overridden, ensuring the correct functioning of hash-based data structures like `HashSet` and `HashMap`.

References